

CSL7070: Computer Architecture

Lecture 3, 15th January 2022

Dip Sankar Banerjee




What is an Instruction Set?

- The complete collection of instructions that are understood by a CPU
- Machine language: binary representation of operations and (addresses of) arguments
- Assembly language: mnemonic representation for humans, e.g.,
OP A,B,C (meaning $A \leftarrow \text{OP}(B,C)$)

Elements of an Instruction

- Operation code (opcode)
 - Do this: ADD, SUB, MPY, DIV, LOAD, STOR
- Source operand reference
 - To this: (address of) argument of op, e.g. register, memory location
- Result operand reference
 - Put the result here (as above)
- Next instruction reference (often implicit)
 - When you have done that, do this: BR

Stack



◆ Implicit operands on stack

◆ Ex. $C = A + B$

Push A


Push B

Add

Pop C

◆ Good code density; used in 60's-70's; now in Java VM

Accumulator



- The accumulator provides an implicit input, and is the implicit place to store the result.
- Ex. $C = A + B$
Load R1, A
Add R3, R1, B
Store R3, c
- Used before 1980

General-purpose Registers

- General-purpose registers are preferred by compilers
 - Reduce memory traffic
 - Improve program speed
 - Improve code density
- Usage of general-purpose registers
 - Holding temporal variables in expression evaluation
 - Passing parameters
 - Holding variables
- GPR and RISC and CISC
 - RISC ISA is extensively used for desktop, server, and embedded: MIPS, PowerPC, UltraSPARC, ARM, MIPS16, Thumb
 - CISC: IBM 360/370, VAX, and Intel 80x86

How Many Registers?

If the number of registers increase:

- ↑ Allocate more variables in registers (fast accesses)
- ↑ Reducing code spill
- ↑ Reducing memory traffic
- ↓ Longer register specifiers (difficult encoding)
- ↓ Increasing register access time (physical registers)
- ↓ More registers to save in context switch

MIPS64: 32 general-purpose registers

Memory Addressing

Instructions see registers, constant values, and memory

- *Addressing mode* decides how to specify an object to access
 - Object can be memory location, register, or a constant
 - Memory addressing is complicated
- *Memory addressing* involves many factors
 - Memory addressing mode
 - Object size
 - byte ordering
 - alignment

For a memory location, its *effective address* is calculated in a certain form of register content, immediate address, and PC, as specified by the addressing mode

Instruction Formats

- When we design ISAs we can look at..
 - Byte-Ordering
 - Instruction Length
 - Number of Opcodes
- ..and their various points and weaknesses..

Byte-Ordering

- How to store data consisting of multiple bytes on a byte-addressable machine?
 - Little Endian
 - **Least significant** byte stored at **lowest** byte address
 - Big Endian
 - **Most significant** byte stored at **lowest** byte address

Address	Example Address	Value
Base + 0	1001	..
Base + 1	1002	..
Base + 2	1003	..
Base +

Byte-Ordering

- Ex. Represent the String

APPLE

		Address				
		Base + 0	Base + 1	Base + 2	Base + 3	Base + 4
Byte-Order	Little Endian	E	L	P	P	A
	Big Endian	A	P	P	L	E

Byte-Ordering


- Little Endian
 - Good for:
 - High-precision arithmetic faster and easier
 - 32 to 16 bit conversion faster (no addition needed)
- Big Endian
 - Good for:
 - Easier to read hex dumps
 - Faster String operations

Byte-Ordering

- Examples of **Little Endian**
 - BMP
 - RTF
 - MS Paint
- Examples of **Big Endian**
 - JPEG
 - Adobe Photoshop
 - MacPaint


Instruction Length

- Fixed Length
 - Ex. MARIE is a fixed length instruction set consisting of 16-bits




- Variable Length

- 12-bits



- 36-bits



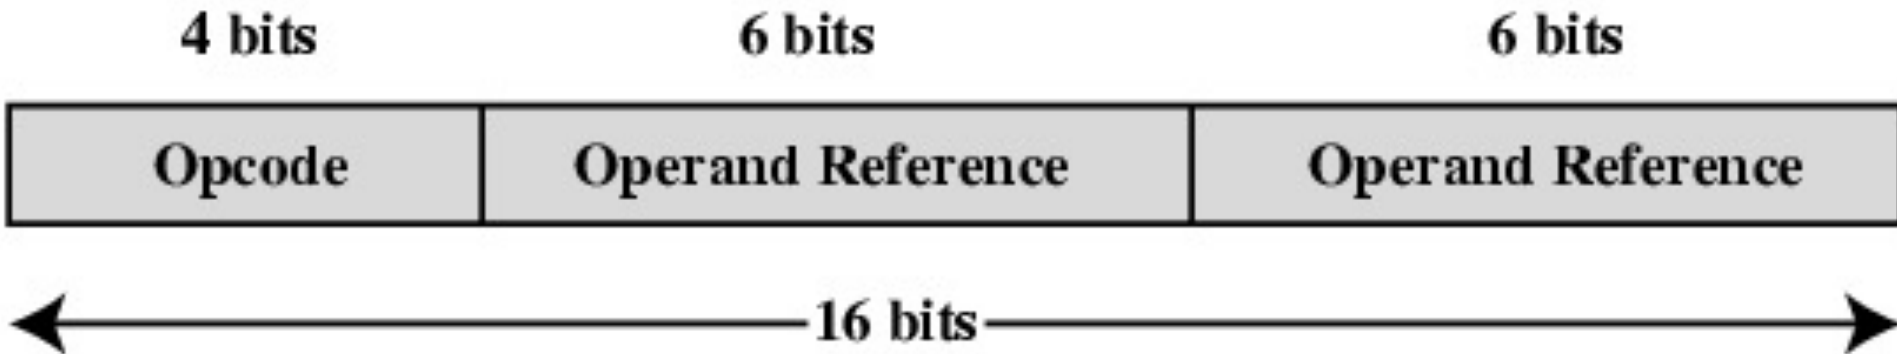
Instruction length

- Fixed Length
 - **Pro**: Decodes faster (Not exactly)
 - Less Complexity
 - **Con**: Wastes Space
 - Opcodes that do not require operands such as MARIE's *halt* makes no use of its address space.
 - Additionally, instructions must be **word aligned**.
Creates **gaps** in memory

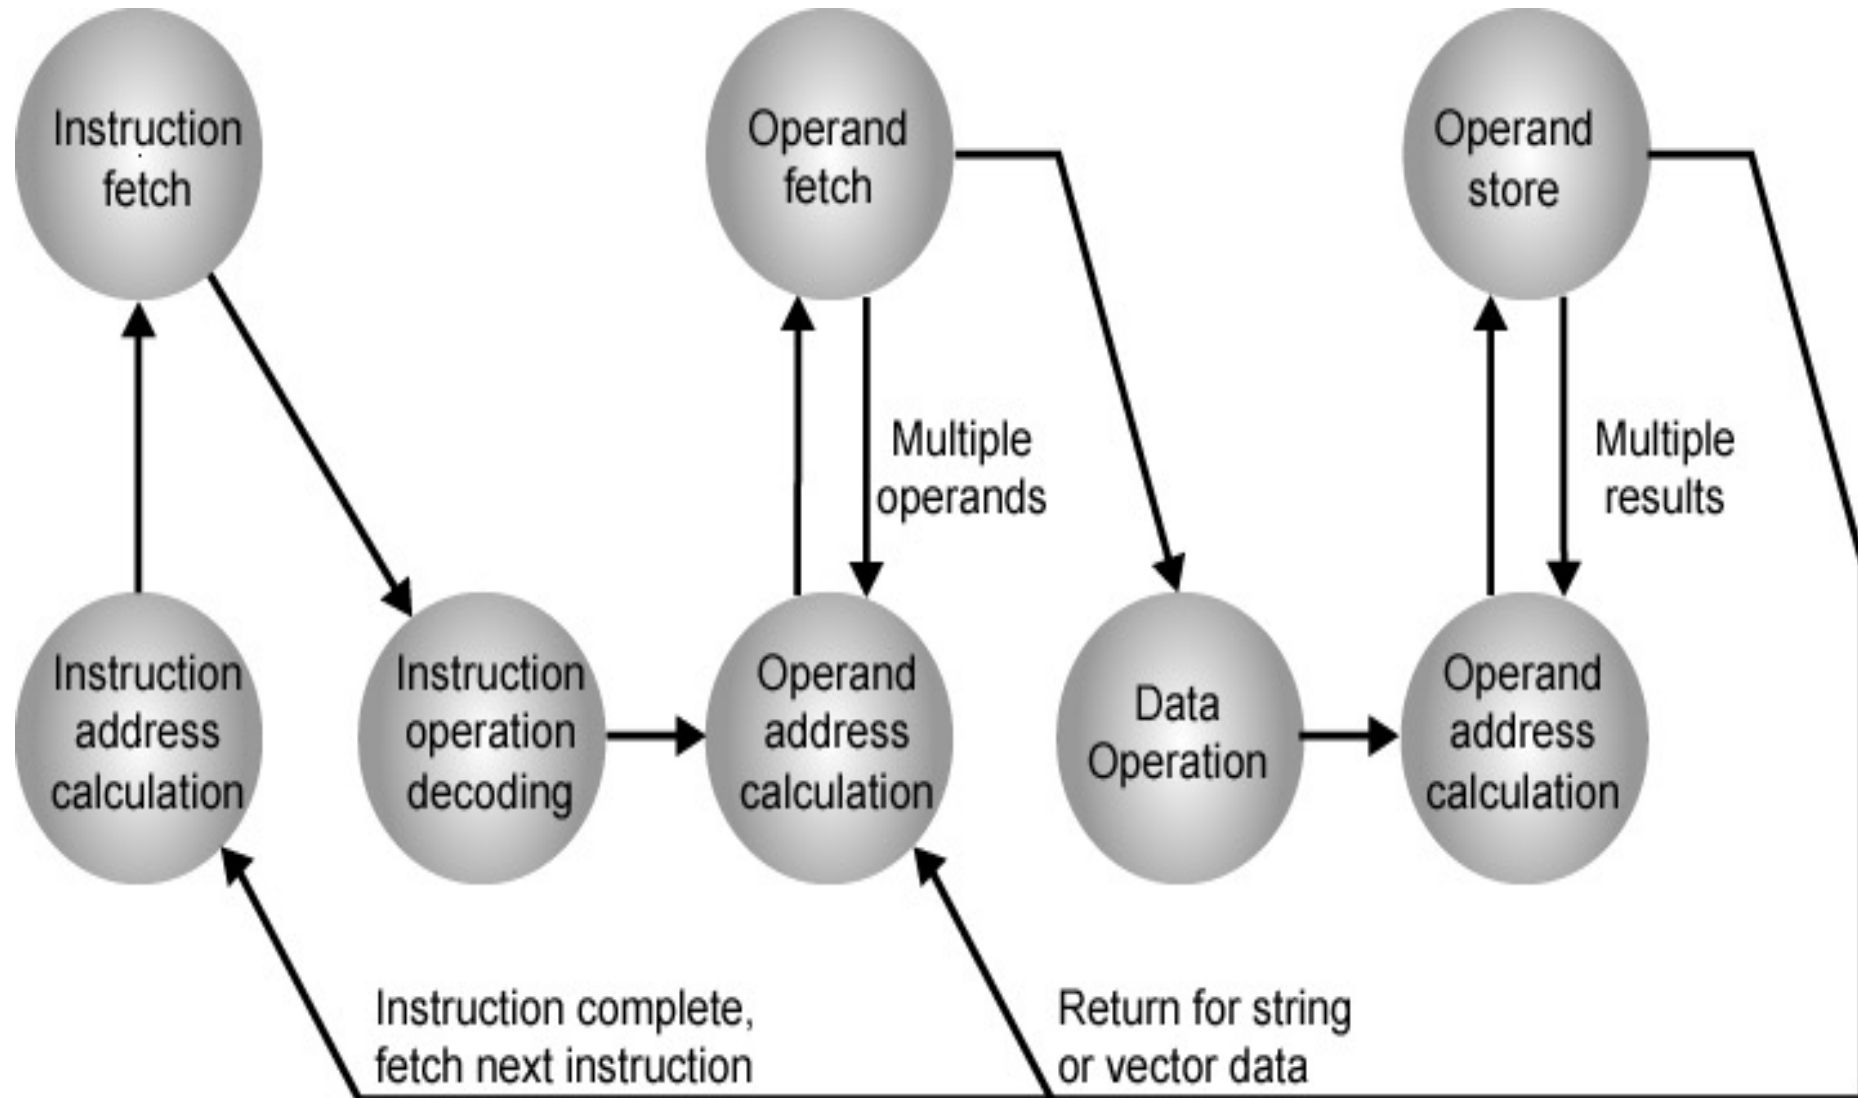
Instruction Length

- Variable Length
 - **Pro**: Saves storage space (Not exactly)
 - Instructions take up only as much space as needed
 - Instructions must be **word aligned** in main memory
 - Therefore instructions of varying lengths will create **gaps** in main memory
 - **Con**: Complex to decode

Simple Instruction Format (using two addresses)



Instruction Cycle State Diagram



Design Decisions (1)

- Operation repertoire
 - How many ops?
 - What can they do?
 - How complex are they?
- Data types (length of words, integer representation)
- Instruction formats
 - Length of op code field
 - Length and number of addresses (e.g., implicit addressing)

Design Decisions (2)

- Registers
 - Number of CPU registers available
 - Which operations can be performed on which registers? General purpose and specific registers
- Addressing modes (see later)
- RISC v CISC

Instruction Types

- Data transfer: registers, main memory, stack or I/O
- Data processing: arithmetic, logical
- Control: systems control, transfer of control

Data Transfer

- Store, load, exchange, move, clear, set, push, pop
- Specifies: source and destination (memory, register, stack), amount of data
- May be different instructions for different (size, location) movements, e.g.,

IBM S/390: L (32 bit word, $R \leftarrow M$), LH (halfword, $R \leftarrow M$), LR (word, $R \leftarrow R$), plus floating-point registers LER, LE, LDR, LD

Or one instruction and different addresses, e.g.
VAX: MOV

Input/Output

- May be specific instructions, e.g. INPUT, OUTPUT
- May be done using data movement instructions (memory mapped I/O)
- May be done by a separate controller (DMA): Start I/O, Test I/O

Arithmetic

- Add, Subtract, Multiply, Divide for signed integer (+ floating point and packed decimal) – may involve data movement
- May include
 - Absolute ($|a|$)
 - Increment ($a++$)
 - Decrement ($a--$)
 - Negate ($-a$)

Logical

- Bitwise operations: AND, OR, NOT, XOR, TEST, CMP, SET
- Shifting and rotating functions, e.g.
 - logical right shift for unpacking: send 8-bit character from 16-bit word
 - arithmetic right shift: division and truncation for odd numbers
 - arithmetic left shift: multiplication without overflow



(a) Logical right shift




(b) Logical left shift




(c) Arithmetic right shift



(d) Arithmetic left shift



(e) Right rotate

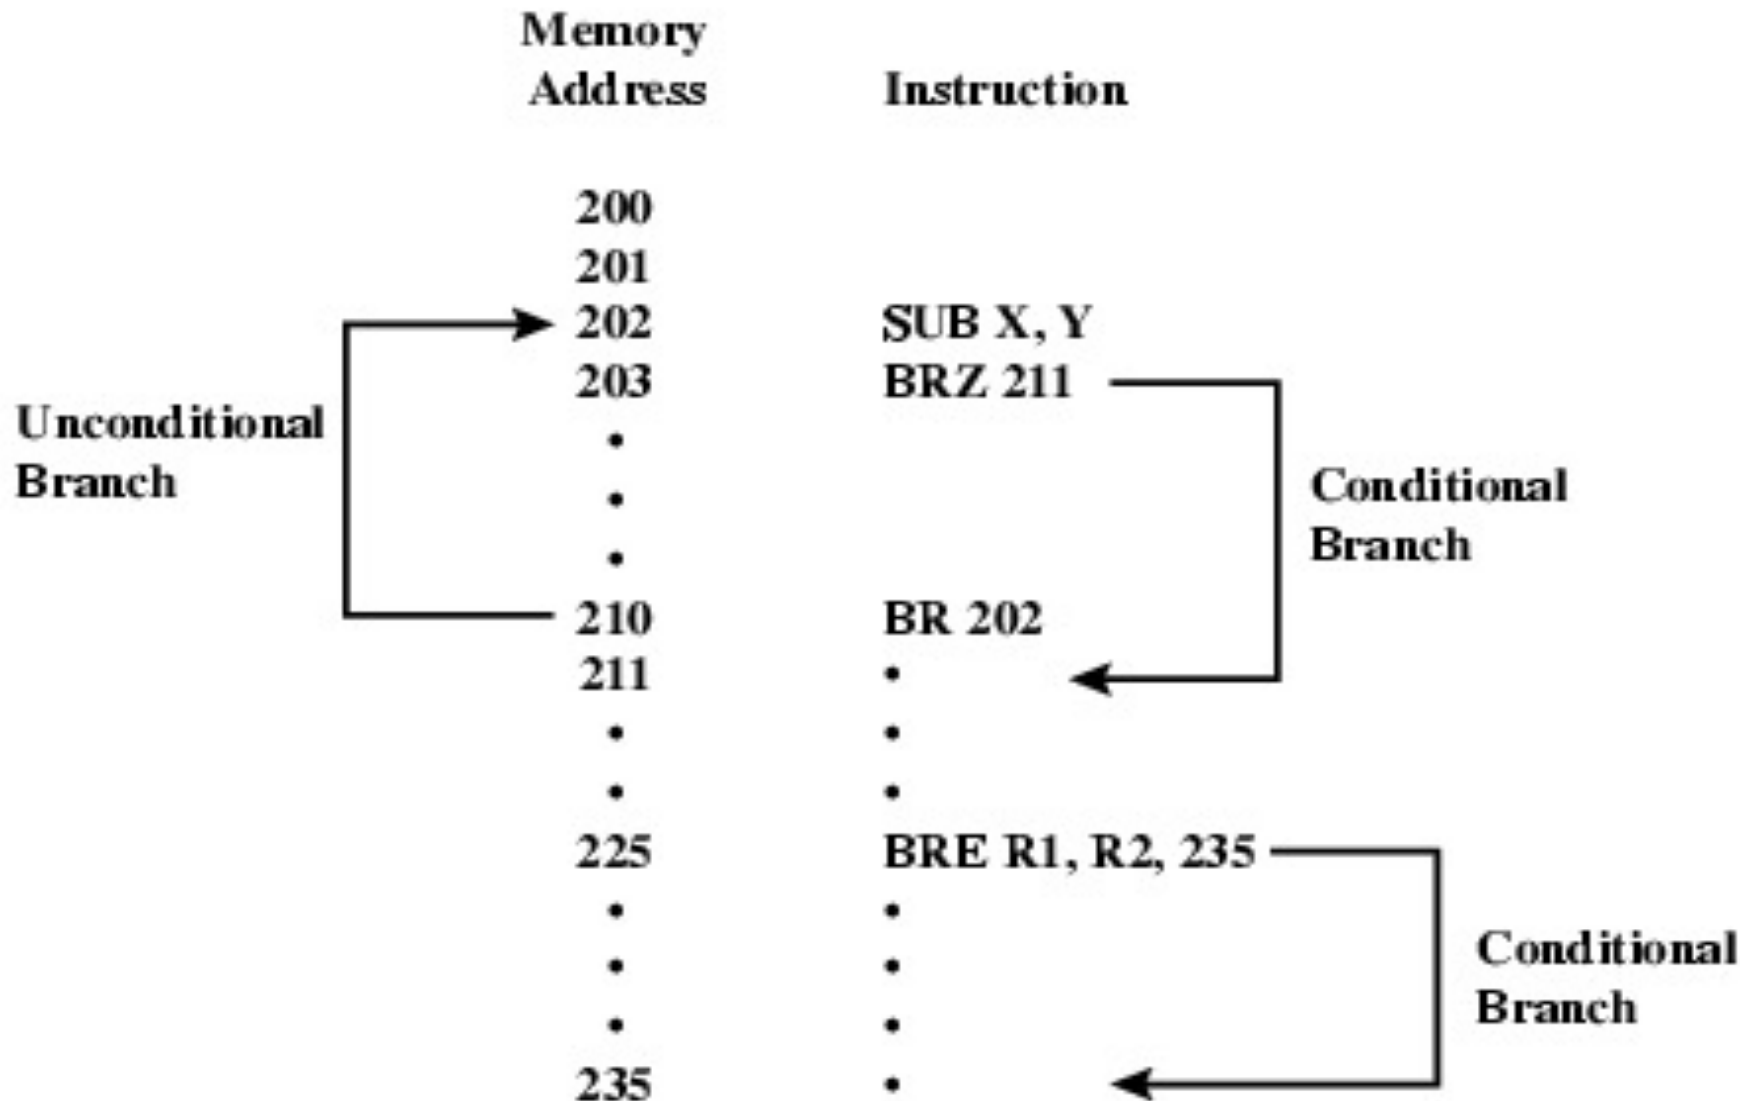


(f) Left rotate

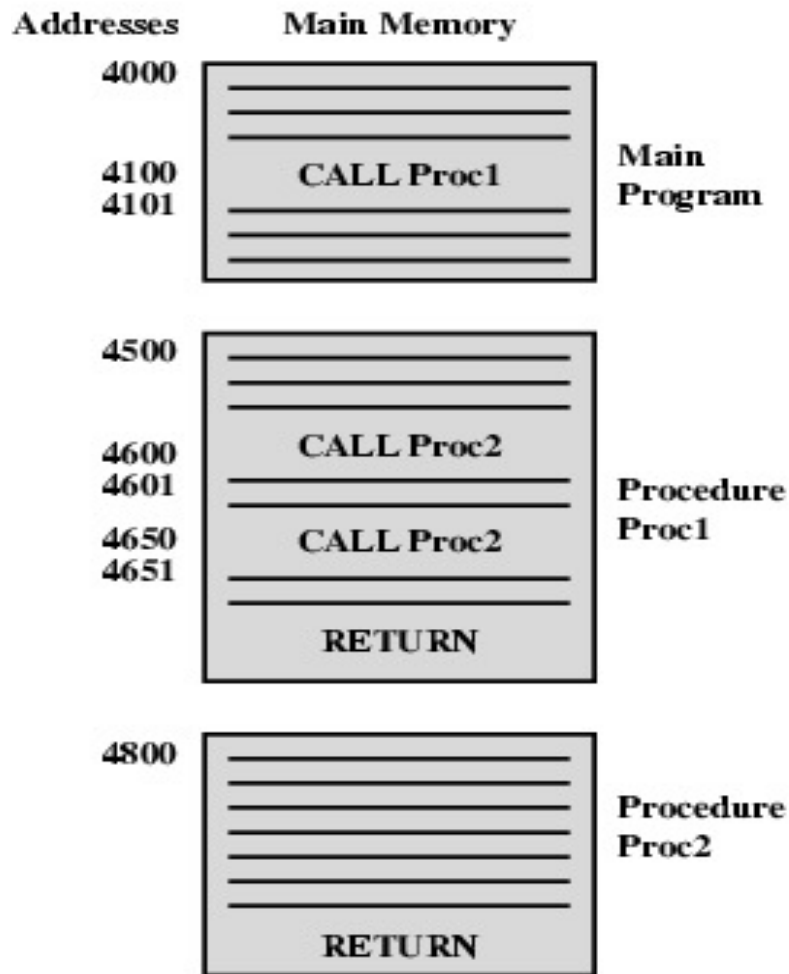
Transfer of Control

- Skip, e.g., increment and skip if zero:
ISZ Reg1, cf. jumping out from loop
- Branch instructions: BRZ X (branch to X if result is zero), BRP X (positive), BRN X (negative), BRE X,R1,R2 (equal)
- Procedure (economy and modularity): call and return


Branch Instruction



Nested Procedure Calls



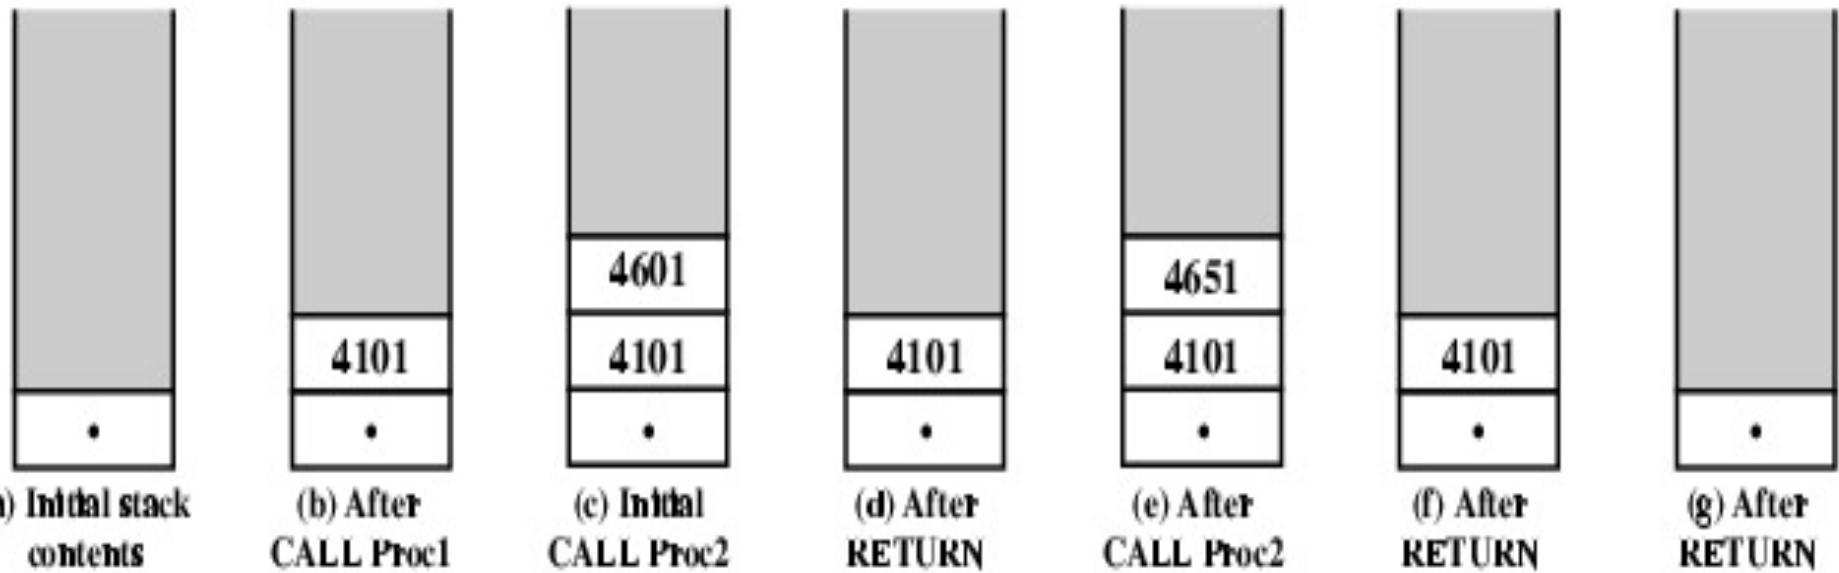
(a) Calls and returns



(b) Execution sequence

Use of Stack

Saving the return address for re-entrant procedures



Types of Operand

- Addresses: immediate, direct, indirect, stack
- Numbers: integer or fixed point (binary, twos complement), floating point (sign, significand, exponent), (packed) decimal (246 = 0000 0010 0100 0110)
- Characters: ASCII (128 printable and control characters + bit for error detection)
- Logical Data: bits or flags, e.g., Boolean 0 and 1

Allocation of Bits

- Number of addressing modes: implicit or additional bits specifying it
- Number of operands
- Register (faster, limited size and number, 32) versus memory
- Number of register sets, e.g., data and address (shorter addresses)
- Address range
- Address granularity (e.g., by byte)

Number of Addresses

- More addresses
 - More complex (powerful?) instructions
 - More registers - inter-register operations are quicker
 - Less instructions per program
- Fewer addresses
 - Less complex (powerful?) instructions
 - More instructions per program, e.g. data movement
 - Faster fetch/execution of instructions
- Example: $Y = (A - B) : [(C + (D \times E))]$

3 addresses

Operation Result, Operand 1, Operand 2

- Not common
- Needs very long words to hold everything

SUB Y,A,B $Y \leftarrow A-B$

MPY T,D,E $T \leftarrow D \times E$

ADD T,T,C $T \leftarrow T+C$

DIV Y,Y,T $Y \leftarrow Y:T$

2 addresses

One address doubles as operand and result

- Reduces length of instruction
- Requires some extra work: temporary storage

MOVE Y,A	$Y \leftarrow A$
SUB Y,B	$Y \leftarrow Y-B$
MOVE T,D	$T \leftarrow D$
MPY T,E	$T \leftarrow T \times E$
ADD T,C	$T \leftarrow T+C$
DIV Y,T	$Y \leftarrow Y:T$

1 address

Implicit second address, usually a register (accumulator, AC)

LOAD D	$AC \leftarrow D$
MPY E	$AC \leftarrow AC \times E$
ADD C	$AC \leftarrow AC + C$
STOR Y	$Y \leftarrow AC$
LOAD A	$AC \leftarrow A$
SUB B	$AC \leftarrow AC - B$
DIV Y	$AC \leftarrow AC : Y$
STOR Y	$Y \leftarrow AC$

0 (zero) addresses

All addresses implicit, e.g. ADD

- Uses a stack, e.g. pop a, pop b, add
- $c = a + b$

Addressing Modes

- Immediate
- Direct
- Indirect
- Register
- Register Indirect
- Displacement (Indexed)
- Stack

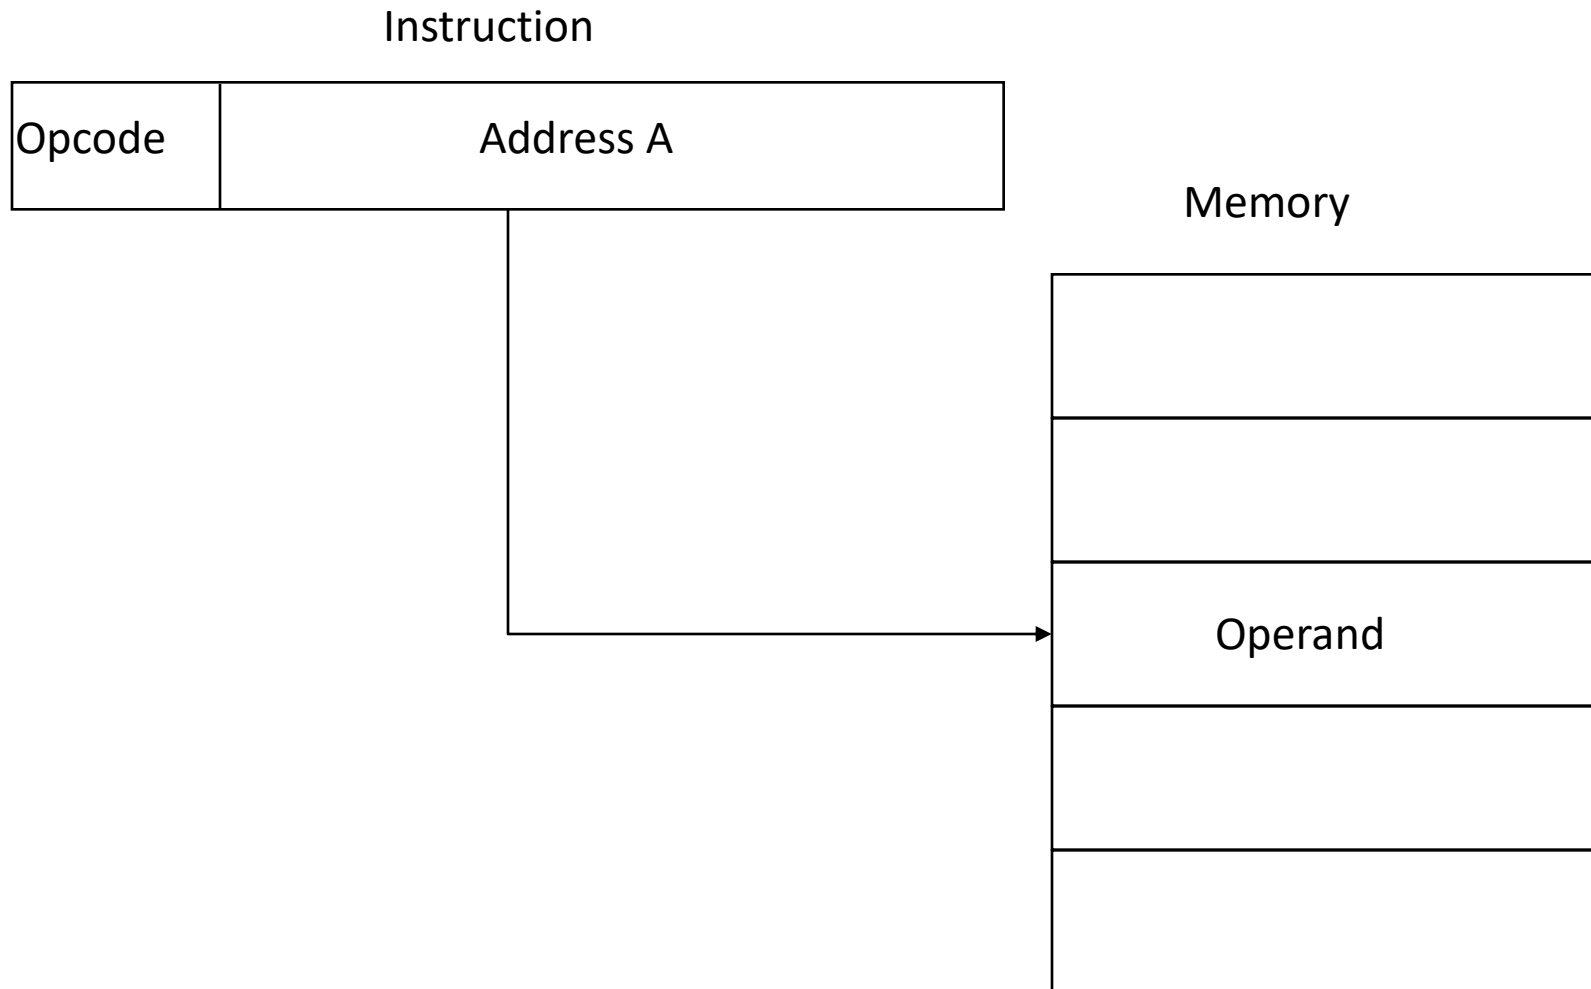
Immediate Addressing

- Operand is part of instruction
- Operand = address field
- e.g., ADD #5
 - Add 5 to contents of accumulator
 - 5 is operand
- No memory reference to fetch data
- Fast
- Limited range

Direct Addressing

- Address field contains address of operand
- Effective address (EA) = address field (A)
- e.g., ADD A
 - Add contents of cell A to accumulator
 - Look in memory at address A for operand
- Single memory reference to access data
- No additional calculations needed to work out effective address
- Limited address space (length of address field)

Direct Addressing Diagram



Indirect Addressing


- Memory cell pointed to by address field contains the address of the operand
- $EA = (A)$
 - Look in A, find effective address and look there for operand
- E.g. ADD (A)
 - Add content of cell pointed to by content of A to accumulator

Indirect Addressing


- Large address space
- 2^n where n = word length
- May be nested, multilevel, cascaded
 - e.g. $EA = (((A)))$
- Multiple memory accesses to find operand
- Hence slower

Indirect Addressing Diagram

Instruction



Memory



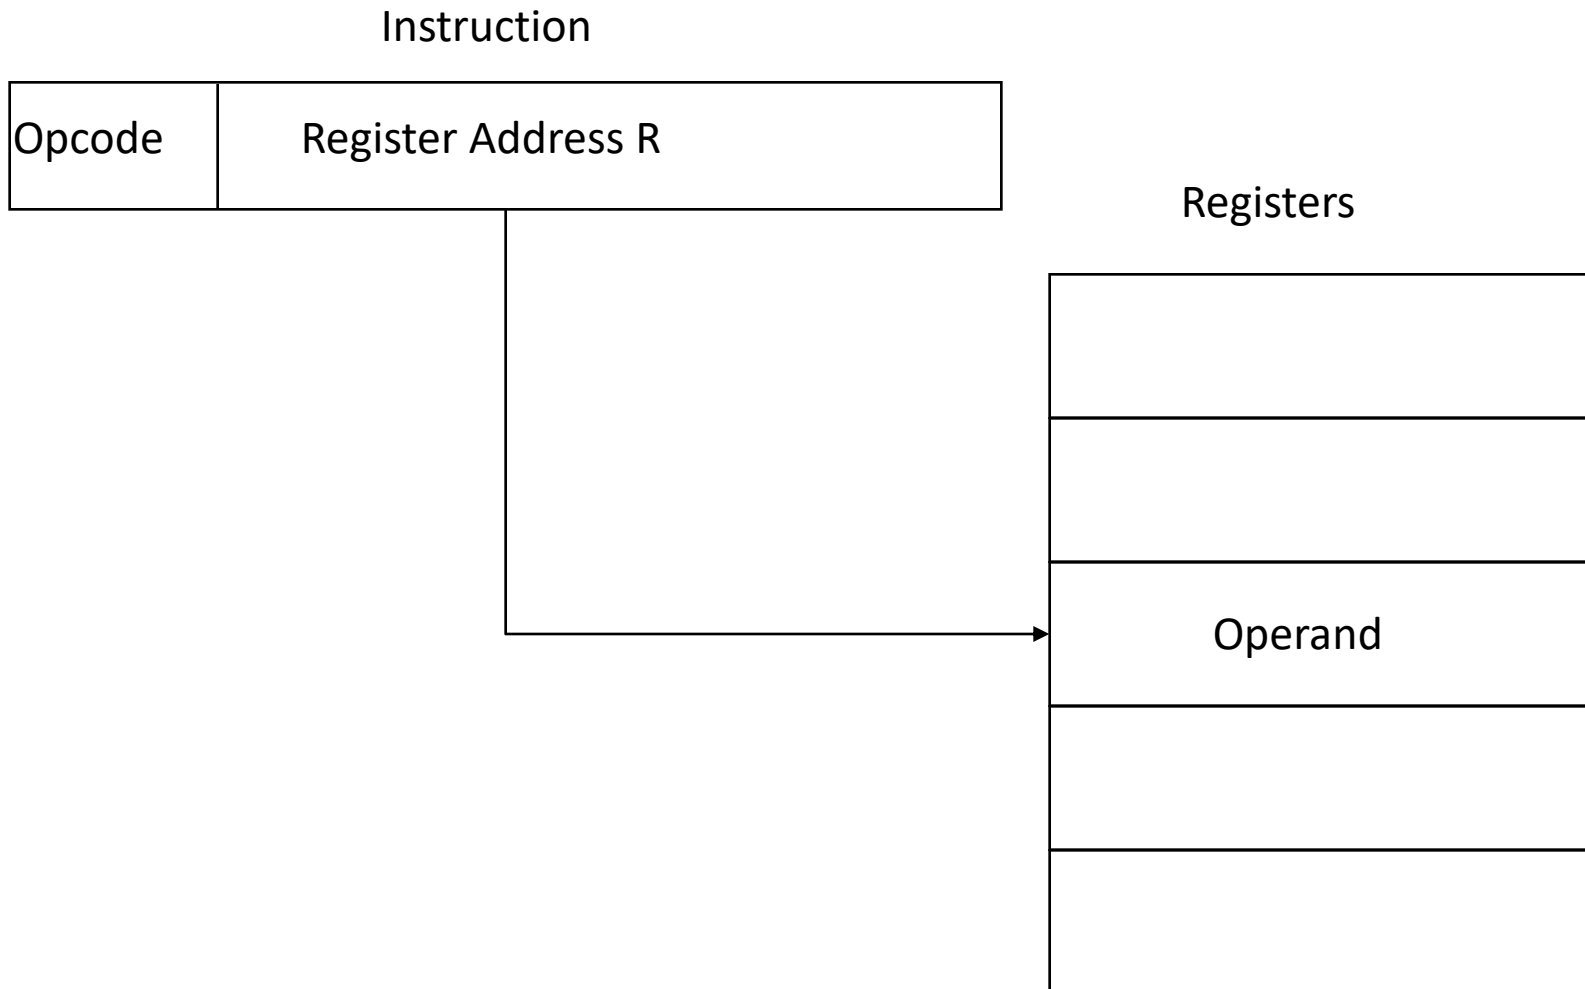
Register Addressing (1)

- Operand is held in register named in address field
- $EA = R$
- Limited number of registers
- Very small address field needed
 - Shorter instructions
 - Faster fetch

Register Addressing (2)

- No memory access
- Very fast execution
- Very limited address space
- Multiple registers helps performance
 - Requires good assembly programming or compiler writing – see register renaming
- cf. direct addressing

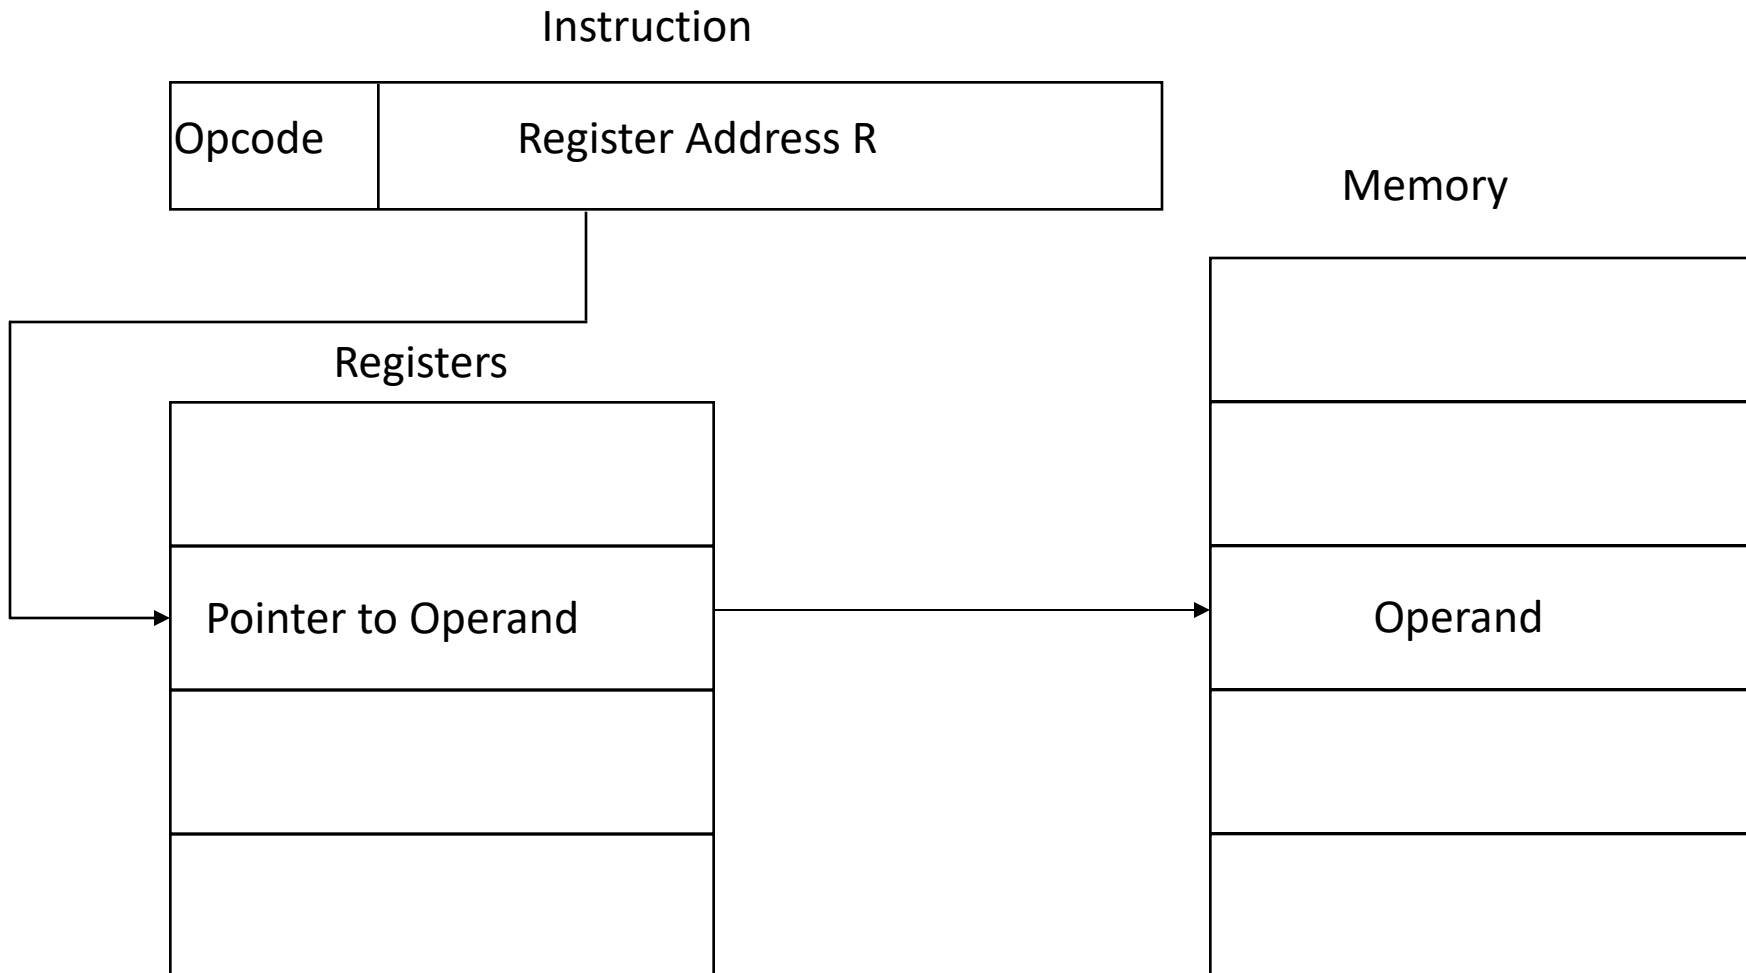
Register Addressing Diagram



Register Indirect Addressing

- Cf. indirect addressing
- $EA = (R)$
- Operand is in memory cell pointed to by contents of register R
- Large address space (2^n)
- One fewer memory access than indirect addressing

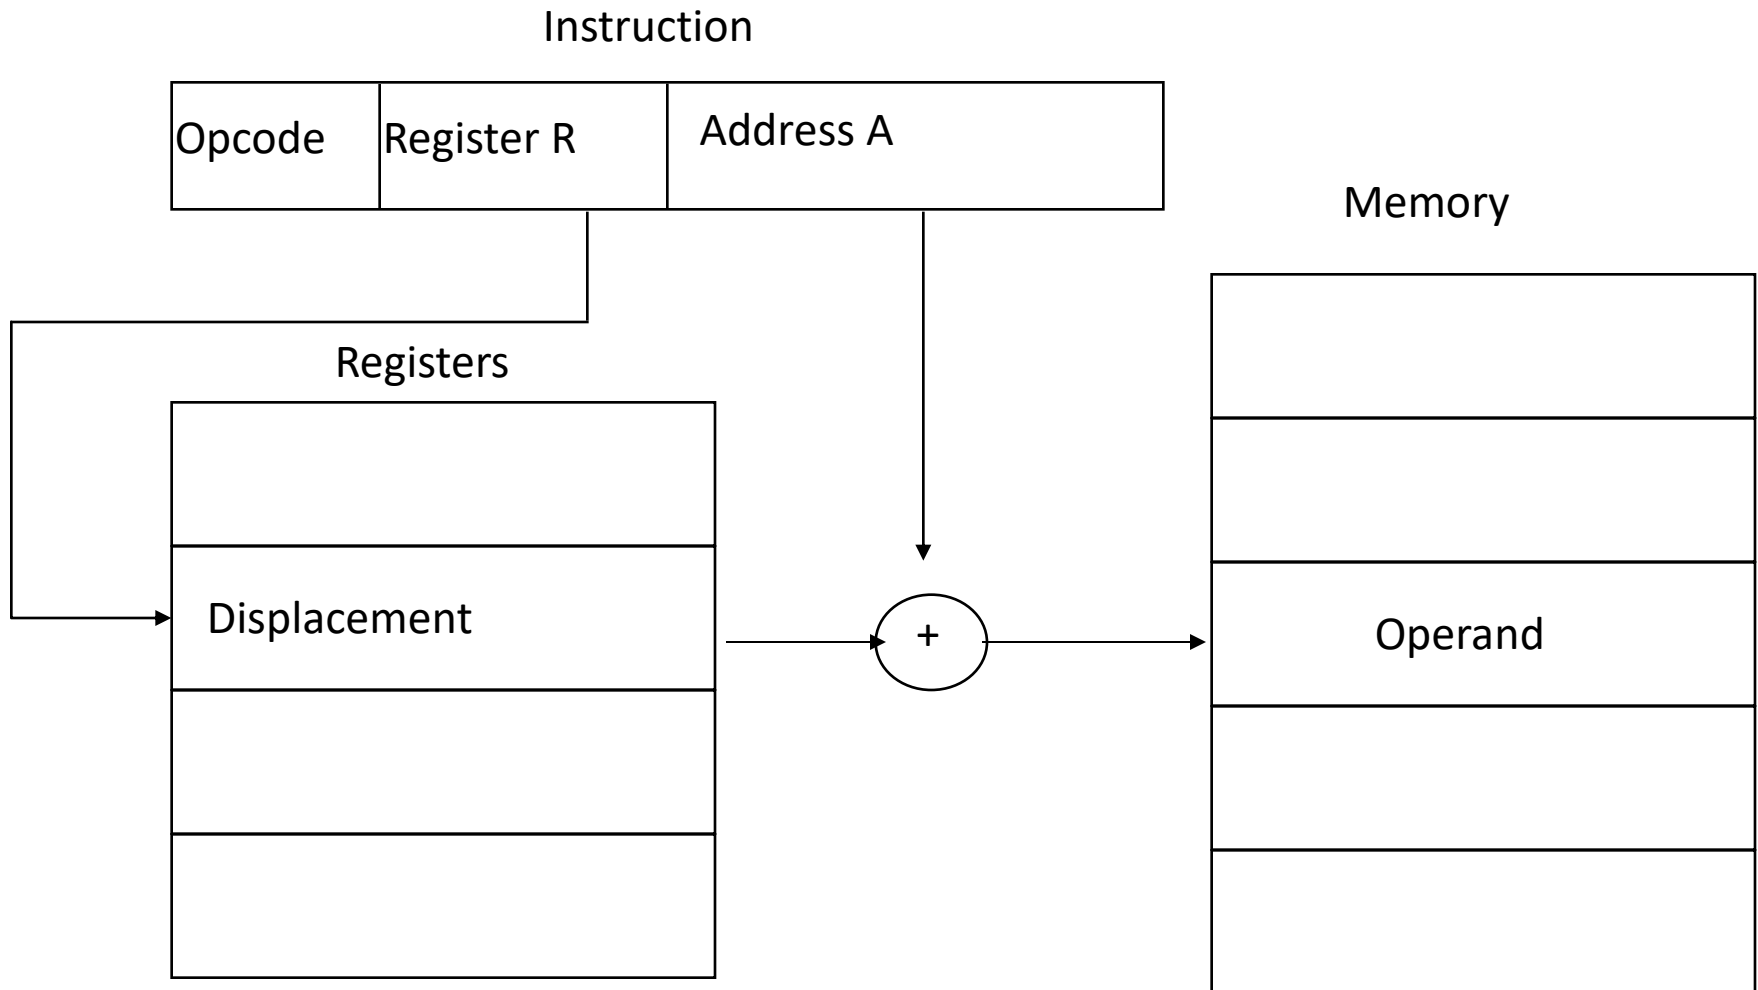
Register Indirect Addressing Diagram



Displacement Addressing

- $EA = A + (R)$
- Address field holds two values
 - A = base value
 - R = register that holds displacement
 - or vice versa
- See segmentation

Displacement Addressing Diagram



Relative/Pseudo-Direct Addressing

- A version of displacement addressing
- $R = \text{Program counter, PC}$
- $EA = A + (PC)$
- i.e., get operand from A cells away from current location pointed to by PC
- cf. locality of reference & cache usage

Indexed Addressing

- $A = \text{base}$
- $R = \text{displacement}$
- $EA = A + R$
- Good for iteration, e.g., accessing arrays
 - $EA = A + R$
 - $R++$
- Sometimes automated: autoindexing (signalled by one bit in instruction)

Stack Addressing

- Operand is (implicitly) on top of stack
- e.g.
 - ADD Pop top two items from stack and add and push result on top