

Introduction to Assembly: RISC-V Instruction Set Architecture

Outline

- Assembly Language
- **RISC-V Architecture**
- Registers vs. Variables
- RISC-V Instructions
- C-to-RISC-V Patterns
- And in Conclusion ...

What is RISC-V?



- Fifth generation of RISC design from UC Berkeley
- A high-quality, license-free, royalty-free RISC ISA specification
 - Implementors do not pay any royalties
 - But see Amdahl's Law:
A decent 180 MHz 32b ARM chip costs \$6 in quantity
A Raspberry Pi (with a 1.8 GHz, quad core ARM and everything else) is \$35:
Licensing cost for the ISA can be in the noise
- Appropriate for all levels of computing system, from micro-controllers to supercomputers
 - 32-bit, 64-bit, and 128-bit variants
 - (we're using 32-bit in class, textbook uses 64-bit)
- Standard maintained by non-profit RISC-V Foundation

Particularly Good For Teaching...

- It is a very very clean RISC
 - No real additional "optimizations"
- Generally only one way to do any particular thing
 - Only exception is two different atomic operation options:
Load Reserved/Store Conditional
Atomic swap/add/etc...
- Clean design for efficient concurrent operations
 - Ground-up understanding of how multiple processors can work together
- Kind to implementers
 - Which means relatively kind when we have you implement one!


Outline

- Assembly Language
- RISC-V Architecture
- **Registers vs. Variables**
- RISC-V Instructions
- C-to-RISC-V Patterns
- And in Conclusion ...

Assembly Variables: Registers

- Unlike HLL like C or Java, assembly does not have variables as you know and love them
 - More primitive, instead what simple CPU hardware can directly support
- Assembly language operands are objects called **registers**
 - **Limited number** of special places to hold values, built directly into the hardware
 - Arithmetic operations can only be performed on these in a RISC!
 - Only memory actions are loads & stores
 - CISC can also perform operations on things **pointed to** by registers
- Benefit:
 - Since registers are directly in hardware, they are very fast to access

Registers live inside the Processor



Speed of Registers vs. Memory

- Given that
 - Registers: 32 words (128 Bytes)
 - Memory (DRAM): Billions of bytes (2 GB to 16 GB on laptop)
- and physics dictates...
 - Smaller is faster
- How much faster are registers than DRAM??
- About 100-500 times faster!
 - in terms of **latency** of one access

Number of RISC-V Registers

- Drawback: Registers are in hardware. To keep them really fast, their number is limited:
 - Solution: RISC-V code must be carefully written to use registers efficiently
- 32 registers in RISC-V, referred to by number x0 – x31
 - Registers are also given symbolic names:
These will be described later and are a "convention"/"ABI" (Application Binary Interface):
Not actually enforced in hardware but needed to follow to keep things consistent
 - Why 32? Smaller is faster, but too small is bad.
 - Plus need to be able to specify 3 registers in operations...
 - Each RISC-V register is 32 bits wide (RV32 variant of RISC-V ISA)
 - Groups of 32 bits called a word in RISC-V ISA
 - P&H CoD textbook uses the 64-bit variant RV64 (explain differences later)
- x0 is special, always holds the value zero
 - So really only 31 registers able to hold variable values

C, Java Variables vs. Registers

- In C (and most HLLs):
 - Variables declared and given a type
 - Example:

```
int fahr, celsius;  
char a, b, c, d, e;
```
 - Each variable can ONLY represent a value of the type it was declared (e.g., cannot mix and match int and char variables)
 - If types are not declared, the object carries around the type with it. EG in python:






```
a = "fubar" # now a is a string  
a = 121 # now a is an integer
```
- In Assembly Language:
 - Registers have ***no type***;
 - Operation determines how register contents are interpreted

RISC-V Memory Alignment...

- RISC-V does not **require** that integers be word aligned...
 - But it is very **very bad** if you don't make sure they are...
- Consequences of unaligned integers
 - Slowdown: The processor is allowed to be a lot slower when it happens
 - In fact, a RISC-V processor may natively only support aligned accesses, and do unaligned-access in **software**!
An unaligned load could take **hundreds of times longer**!
 - Lack of **atomicity**: The whole thing doesn't happen at once...
can introduce lots of very subtle bugs
- So in **practice**, RISC-V requires integers to be aligned on 4-byte boundaries

RISC-V Instructions

- Instructions are fixed, 32b long
 - Must also be word aligned, or half-word aligned if the 16b optional (C) instruction set is also enabled
- Only a few formats (we'll go into detail later)...




Outline

- Assembly Language
- RISC-V Architecture
- Registers vs. Variables
- **RISC-V Instructions**
- C-to-RISC-V Patterns
- And in Conclusion ...

RISC-V Instruction Assembly Syntax

- Instructions have an opcode and operands

E.g., **add x1, x2, x3 # x1 = x2 + x3**



Addition and Subtraction of Integers

- Addition in Assembly

- Example: `add x1, x2, x3` (in RISC-V)
- Equivalent to: `a = b + c` (in C)

where C variables \Leftrightarrow RISC-V registers are:

`a \Leftrightarrow x1, b \Leftrightarrow x2, c \Leftrightarrow x3`

- Subtraction in Assembly

- Example: `sub x3, x4, x5` (in RISC-V)
- Equivalent to: `d = e - f` (in C)

where C variables \Leftrightarrow RISC-V registers are:

`d \Leftrightarrow x3, e \Leftrightarrow x4, f \Leftrightarrow x5`

No-Op

- A No-op is an instruction that does nothing...
 - Why?
You may need to replace code later: No-ops can fill space, align data, and perform other options
- By **convention** RISC-V has a specific no-op instruction...
 - `add x0 x0 x0`
- Why?
 - Writes to x0 are always ignored...
RISC-V uses that a lot as we will see in the jump-and-link operations
 - Making a "standard" no-op improves the disassembler and can potentially improve the processor
 - Special case the particular conventional no-op.

Addition and Subtraction of Integers

Example 1

- How to do the following C statement?

a = b + c + d - e;

- Break into multiple instructions

```
add x1, x2, x3    # temp = b + c
add x1, x1, x4    # temp = temp + d
sub x1, x1, x5    # a = temp - e
```

- A single line of C may turn into several RISC-V instructions

add x3, x4, x0 (in RISC-V) same
f = g (in C)

Immediates

- *Immediates are used to provide numerical constants*
- Constants appear often in code, so there are special instructions for them:

- Ex: Add Immediate:

addi x3,x4,-10 (in RISC-V)

f = g - 10 (in C)

where RISC-V registers **x3, x4** are associated with C variables **f, g**

- Syntax similar to `add` instruction, except that last argument is a number instead of a register

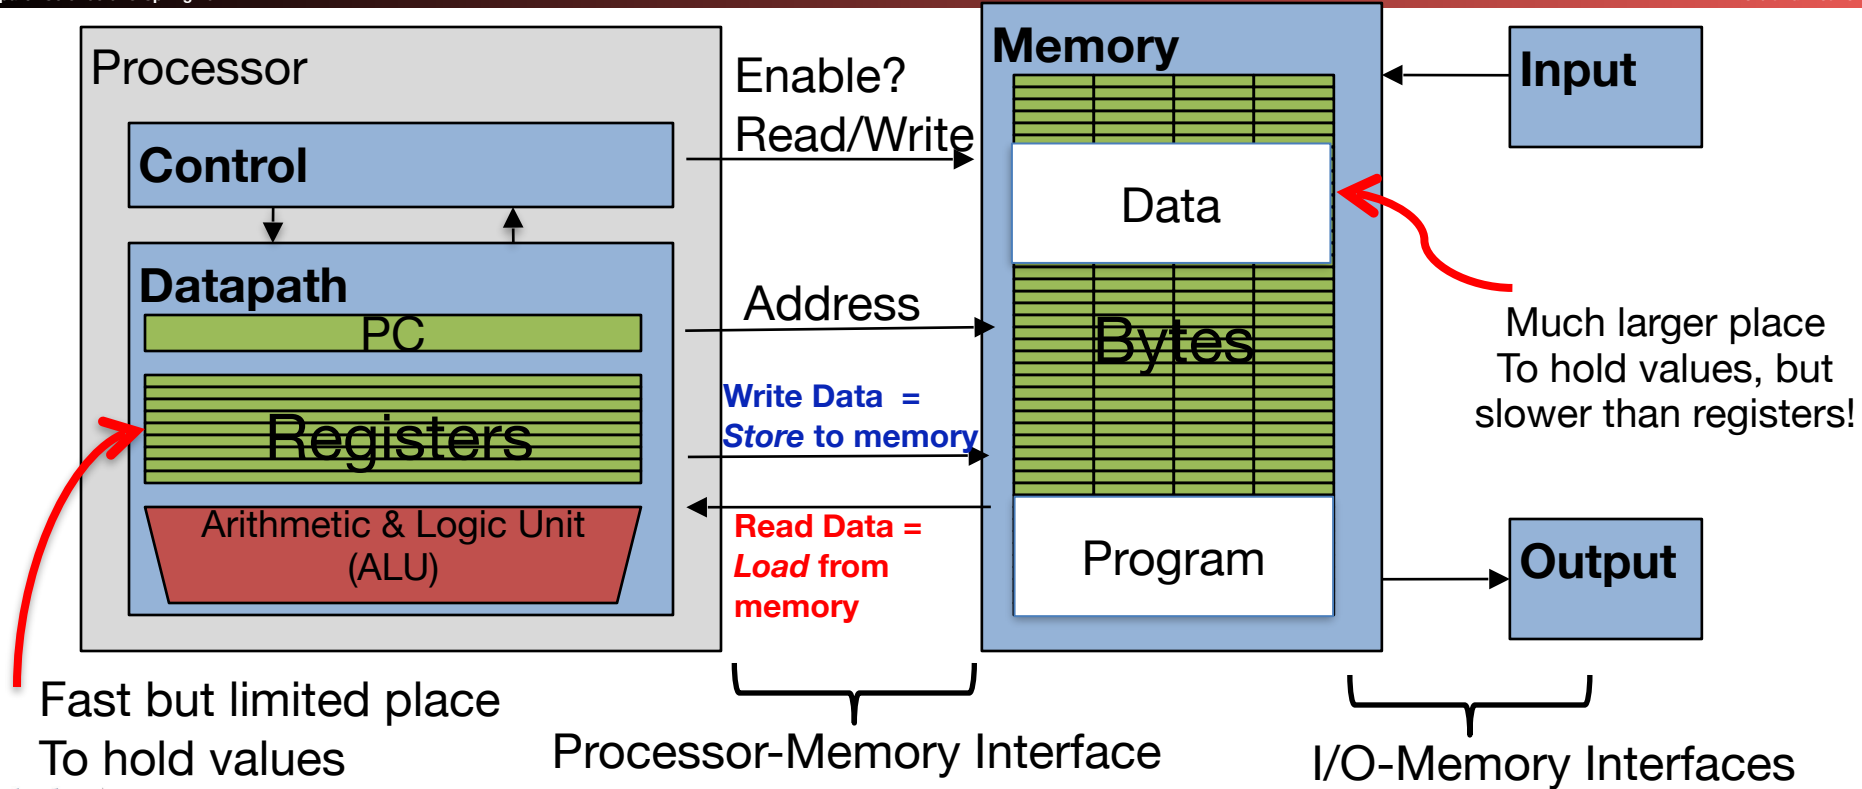
addi x3,x4,0 (in RISC-V) same as
f = g (in C)

Immediates & Sign Extension...

- Immediates are necessarily small
 - An I-type instruction can only have 12 bits of immediate
- In RISC-V immediates are "sign extended"
 - So the upper bits are the same as the largest bit
- So for a 12b immediate...
 - Bits 31:12 get the same value as Bit 11

Data Transfer:


Load from and Store to memory



Memory Addresses are in Bytes

- Data typically smaller than 32 bits, but rarely smaller than 8 bits (e.g., char type)
 - So everything is a multiple of 8 bits
- Remember, 8 bit chunk is called a byte (1 word = 4 bytes)
- Memory addresses are really in bytes, not words
- Word addresses are 4 bytes apart
 - Word address is same as address of rightmost byte – least-significant byte (i.e. Little-endian convention)

Least-significant byte in word



Least-significant byte gets the smallest address

Transfer from Memory to Register

- C code

```
int A[100];  
g = h + A[3];
```
- Using Load Word (**lw**) in RISC-V:

```
lw  x10,12(x13) # Reg x10 gets A[3]  
add x11,x12,x10 # g = h + A[3]
```

Assume: x13 – base register (pointer to A[0])

Note: 12 – offset in bytes

Offset must be a constant known at *assembly time*

Transfer from Register to Memory

- C code

```
int  A[100];  
A[10] = h + A[3];
```
- Using Store Word (sw) in RISC-V:

```
lw  x10,12(x13)  # Temp reg x10 gets A[3]  
add x10,x12,x10  # Temp reg x10 gets h + A[3]  
sw  x10,40(x13)  # A[10] = h + A[3]
```

Assume: $x13$ – base register (pointer)

Note: $12, 40$ – offsets in bytes

$x13+12$ and $x13+40$ must be multiples of 4 to maintain alignment

Loading and Storing Bytes

- In addition to word data transfers (lw, sw), RISC-V has **byte** data transfers:
 - load byte: **lb**
 - store byte: **sb**
 - Same format as lw, sw
 - E.g., **lb x10, 3(x11)**
 - contents of memory location with address = sum of “3” + contents of register x11 is copied to the low byte position of register **x10**.
- RISC-V also has “unsigned byte” loads (**lbui**) which zero extend to fill register. Why no unsigned store byte **sbui**?

sfers:

RISC-V also has “unsigned byte” loads (**lbu**) which zero extend to fill register. Why no unsigned store byte **sbu**?

x10: xxxxxx xxxxxx xxxxxx xxxxxx xxxxxx xxxxxx

...is copied to "sign-extend"

This bit

byte loaded

Example - Tracing Assembly Code

```
addi x11,x0,0x3f5
```

```
sw x11,0(x5)
```

```
lb x12,1(x5)
```

What's the value in x12?

Answer	x12
A	0x5
B	0xf
C	0x3
D	0xffffffff

Example - Tracing Assembly Code

```
addi x11,x0,0x3f5
```

```
sw x11,0(x5)
```

```
lb x12,1(x5)
```

What's the value in x12?

Answer	x12
A	0x5
B	0xf
C	0x3
D	0xffffffff

Note Endianness...

- Remember, RISC-V is "little endian"
 - `byte[0]` = least significant byte of the number
 - `byte[3]` = most significant byte of the number
- So for this example...
 - `byte[0]` = `0xf5`
 - `byte[1]` = `0x03`
 - `byte[2]` = `0x00`
 - `byte[3]` = `0x00`

Another Example

```
addi x11,x0,0x8f5
```

```
sw x11,0(x5)
```

```
lb x12,1(x5)
```

What's the value in x12?

Answer	x12
A	0x8
B	0xf8
C	0x3
D	0xffffffff8

Example - Tracing Assembly Code

```
addi x11,x0,0x8f5
```

```
sw x11,0(x5)
```

```
lb x12,1(x5)
```

What's the value in x12?

Answer	x12
A	0x8
B	0xf8
C	0x3
D	0xffffffff8

Two Reasons for The Answer...

- The immediate got sign extended...
 - So `0xffffffff8f5` got written
- Then load byte is called
 - So it will load `byte[1]`, which is `0xf8`
- But load byte sign extends too...
 - So what gets loaded into the register is `0xfffffffff8`
- If we did `lbu` we'd instead get `0xf8`

RISC-V Logical Instructions

Useful to operate on fields of bits within a word

e.g., characters within a word (8 bits)

Operations to pack /unpack bits into words

Called logical operations

Logical operations	C operators	Java operators	RISC-V instructions
Bit-by-bit AND	&	&	and
Bit-by-bit OR			or
Bit-by-bit XOR	^	^	xor
Shift left logical	<<	<<	sll
Shift right	>>	>>	srl/sra

Logical Shifting

- Shift Left Logical: `slli x11, x12, 2` # $x11 = x12 \ll 2$
 - Store in x11 the value from x12 shifted 2 bits to the left (they fall off end), inserting 0's on right; \ll in C

Before: `0000 0002`_{hex}

`0000 0000 0000 0000 0000 0000 0000 0010`_{two}

After: `0000 0008`_{hex}

`0000 0000 0000 0000 0000 0000 0000 1000`_{two}

What arithmetic effect does shift left have?

- Shift Right Logical: `srl` is opposite shift; \gg
 - Zero bits inserted at left of word, right bits shifted off end

Arithmetic Shifting

- *Shift right arithmetic* (**srai**) moves n bits to the right (insert high-order sign bit into empty bits)
- For example, if register x10 contained
 $1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1110\ 0111_{\text{two}} = -25_{\text{ten}}$
- If execute `sra x10, x10, 4`, result is:
 $1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1110_{\text{two}} = -2_{\text{ten}}$
- Unfortunately, this is NOT same as dividing by 2^n
 - Fails for odd negative numbers
 - C arithmetic semantics is that division should round towards 0

Transfer Array Value from Memory to Register With Variable Indexing

- C code

```
int A[100]; /* x13 */  
int i;      /* x14 */
```


```
...  
g = h + A[3]; /* h = x12, g = x11, tmp = x15 */
```

- Using Load Word (`lw`) in RISC-V with pointer arithmetic:

```
sll x15,x14,2 /* Multiply by 4 for ints */  
add x15,x15,x13 /* A + 4 * i */  
lw x10,0(x15)  
add x11,x12,x10
```

Computer Decision Making

- Based on computation, do something different
- Normal operation on CPU is to execute instructions in sequence
- Need special instructions for programming languages: *if*-statement
- RISC-V: *if*-statement instruction is
beq register1, register2, L1
means: go to instruction labeled L1
if (value in register1) == (value in register2)
....otherwise, go to next instruction
- **beq** stands for *branch if equal*
- Other instruction: **bne** for *branch if not equal*



Types of Branches

- **Branch** – change of control flow
- **Conditional Branch** – change control flow depending on outcome of comparison
 - branch *if* equal (**b_{eq}**) or branch *if not* equal (**b_{ne}**)
 - Also branch if less than (**b_{lt}**) and branch if greater than or equal (**b_{ge}**)
- **Unconditional Branch** – always branch
 - a RISC-V instruction for this: *jump (j)*
 - *We will see later that j doesn't exist (it's a "pseudo-instruction")*

Outline

- Assembly Language
- RISC-V Architecture
- Registers vs. Variables
- RISC-V Instructions
- C-to-RISC-V Patterns
- And in Conclusion ...

Labels In Assembly Language...

- We commonly see "labels" in the code
 - `foo: add x2 x1 x0`
- The assembler converts these into positions in the code
 - What address in the code that label is at...
- Then when you have control flow instructions like jumps and branches...
 - e.g. `bne x0 x2 foo`
- The assembler in outputting the code does the necessary math so the jump or branch will go to the right place

Example *if* Statement

- Assuming assignments below, compile *if* block

$f \rightarrow x10$ $g \rightarrow x11$ $h \rightarrow x12$

$i \rightarrow x13$ $j \rightarrow x14$

```
if (i == j)
    f = g + h;
```

```
bne x13,x14,done
```

```
add x10,x11,x12
```

```
done:
```

Example *if-else* Statement

- Assuming assignments below, compile

$f \rightarrow x10$ $g \rightarrow x11$ $h \rightarrow x12$ $i \rightarrow x13$ $j \rightarrow x14$

if ($i == j$)	bne x13,x14,else
$f = g + h;$	add x10,x11,x12
else	j done
$f = g - h;$	else: sub x10,x11,x12
	done:

Magnitude Compares in RISC-V

- Until now, we've only tested equalities (`==` and `!=` in C); General programs need to test `<` and `>` as well.
- RISC-V magnitude-compare branches:

“Branch on Less Than”

Syntax: `blt reg1, reg2, label`

Meaning: `if (reg1 < reg2) // Registers are signed
goto label;`

- “Branch on Less Than Unsigned”

Syntax: `bltu reg1, reg2, label`

Meaning: `if (reg1 < reg2) // treat registers as unsigned integers
goto label;`

Berkeley EECS “Branch on Greater Than or Equal” (and its unsigned version) also exists.

But RISC philosophy...

- A CISC might also have "branch if greater than"...
 - But RISC-V doesn't.
- Instead you can switch the argument
 - branch if greater then reg1 reg2...
 - branch if less than reg2 reg1
- So we have pseudo-instructions, where the assembler converts things
 - `bgt x2 x3 foo` becomes
 - `blt x3 x2 foo`

C Loop Mapped to RISC-V Assembly

```
int A[20];
int sum = 0;
for (int i=0; i<20; i++)
    sum += A[i];
```

```
# Assume x8 holds pointer to A
# Assign x10=sum, x11=i
add x10, x0, x0 # sum=0
add x11, x0, x0 # i=0
addi x12,x0,20  # x12=20
Loop:
bge x11, x12, exit:
sll x13, x11, 2    # i * 4
add x13, x13, x8   # & of A + i
lw x13, 0(x13)     # *(A + i)
add x10, x10, x13  # increment sum
addi x11, x11, 1   # i++
j Loop             # Iterate
exit:
```

Comments...

- The simple translation is suboptimal!
 - A more efficient way:
- Inner loop is now 4 instructions rather than 7
 - And only 1 branch/jump rather than two:
Because first time through is always true so can move check to the end!
- The compiler will often do this automatically for optimization
 - See that i is only used as an index in a loop

```
# Assume x8 holds pointer to A
# Assign x10=sum
add  x10, x0, x0    # sum=0
add  x11, x0, x8    # Copy of A
addi x12, x11, 80   # x12=80 + A
loop:
lw    x13, 0(x11)
add   x10, x10, x13
addi  x11, x11, 4
blt   x11, x12, loop
```

And Premature Optimization...

- In general we want **correct** translations of C to RISC-V
- It is **not** necessary to optimize
 - Just translate each C statement on its own
- Why?
 - Correctness first, performance second
 - Getting the wrong answer fast is not what we want from you...
 - We're going to need to read your assembly to grade it!
 - Multiple ways to optimize, but the straightforward translation is mostly unique-ish.

Outline

- Assembly Language
- RISC-V Architecture
- Registers vs. Variables
- RISC-V Instructions
- C-to-RISC-V Patterns
- And in Conclusion ...

In Conclusion,...

- Instruction set architecture (ISA) specifies the set of commands (instructions) a computer can execute
- Hardware registers provide a few very fast variables for instructions to operate on
- RISC-V ISA requires software to break complex operations into a string of simple instructions, but enables faster, simple hardware
- Assembly code is human-readable version of computer's native machine code, converted to binary by an *assembler*