


CSL7070: Computer Architecture

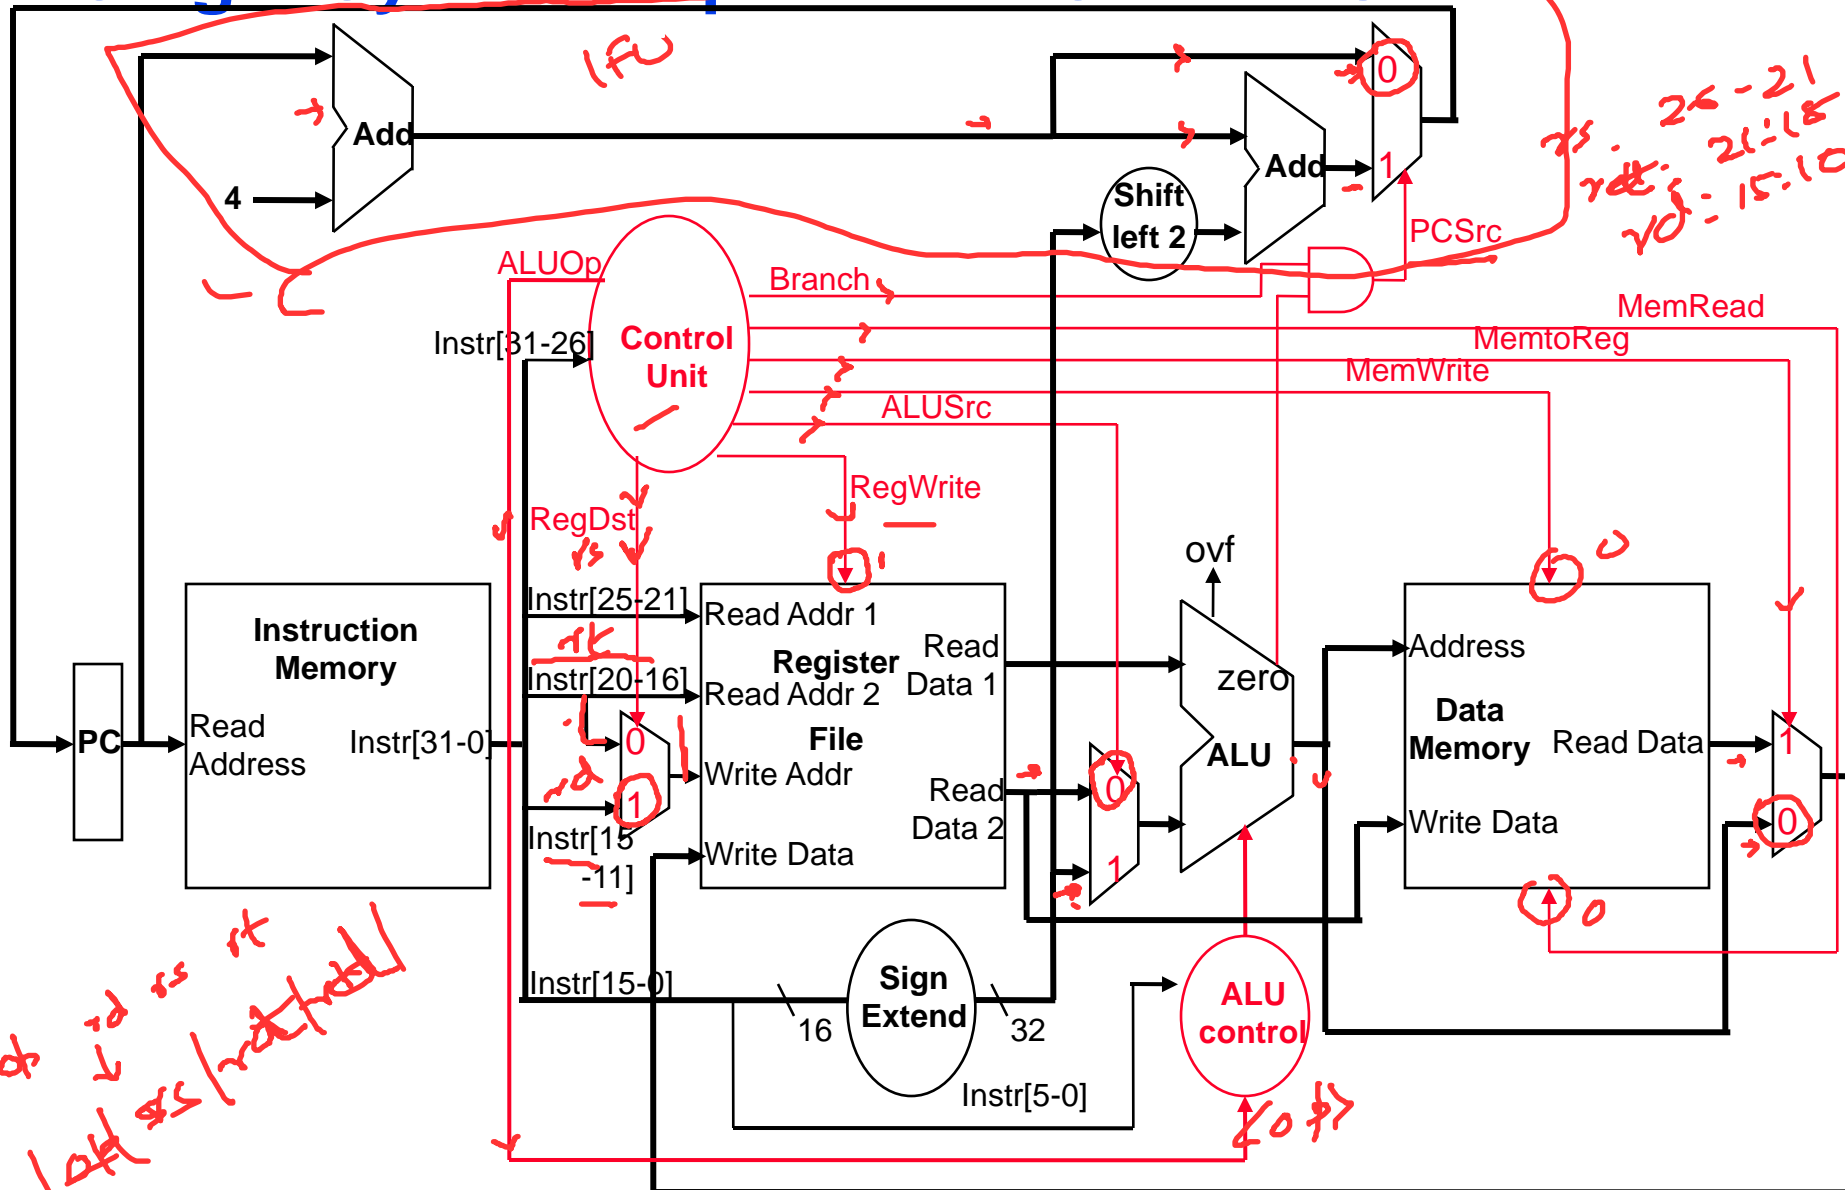
Lecture 9, 21st February 2022

Dip Sankar Banerjee

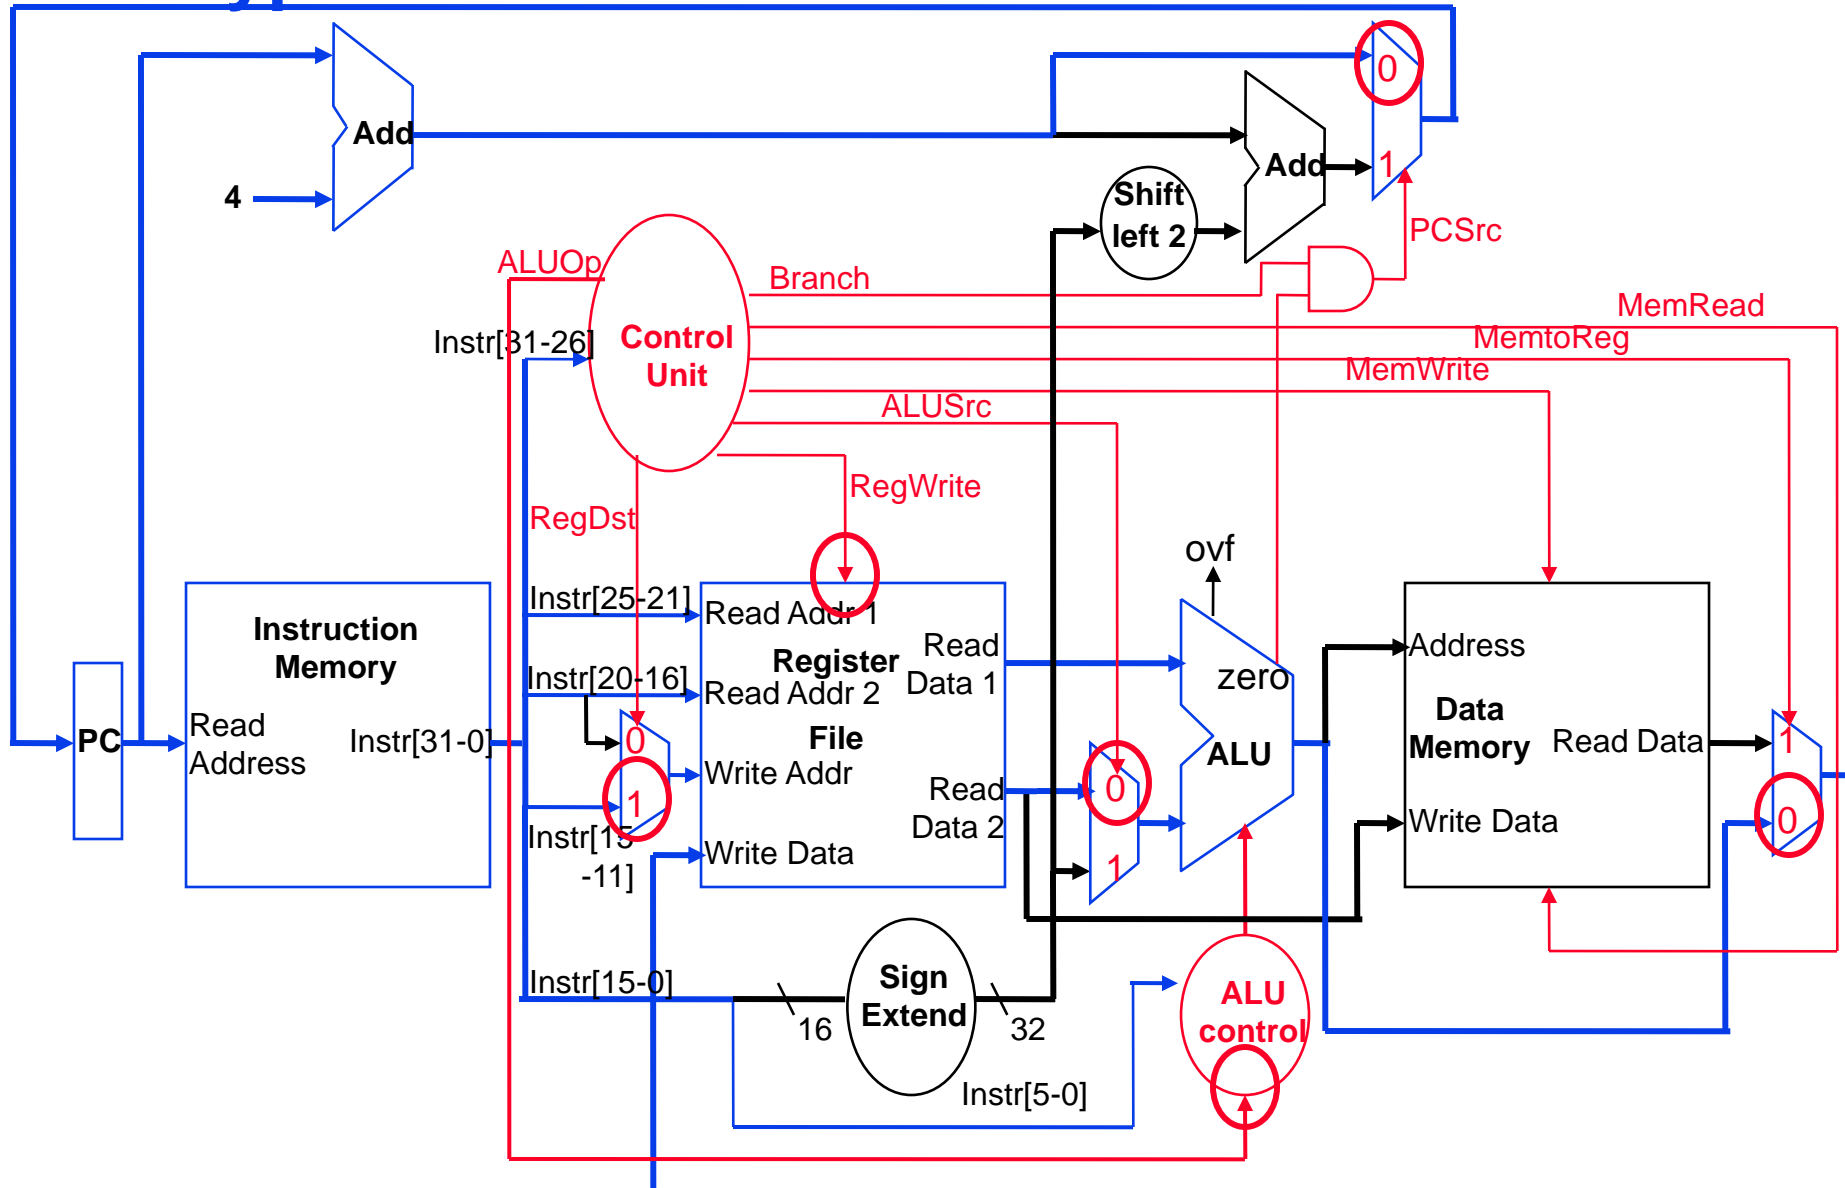


Indian Institute of Technology, Jodhpur
January-April 2022

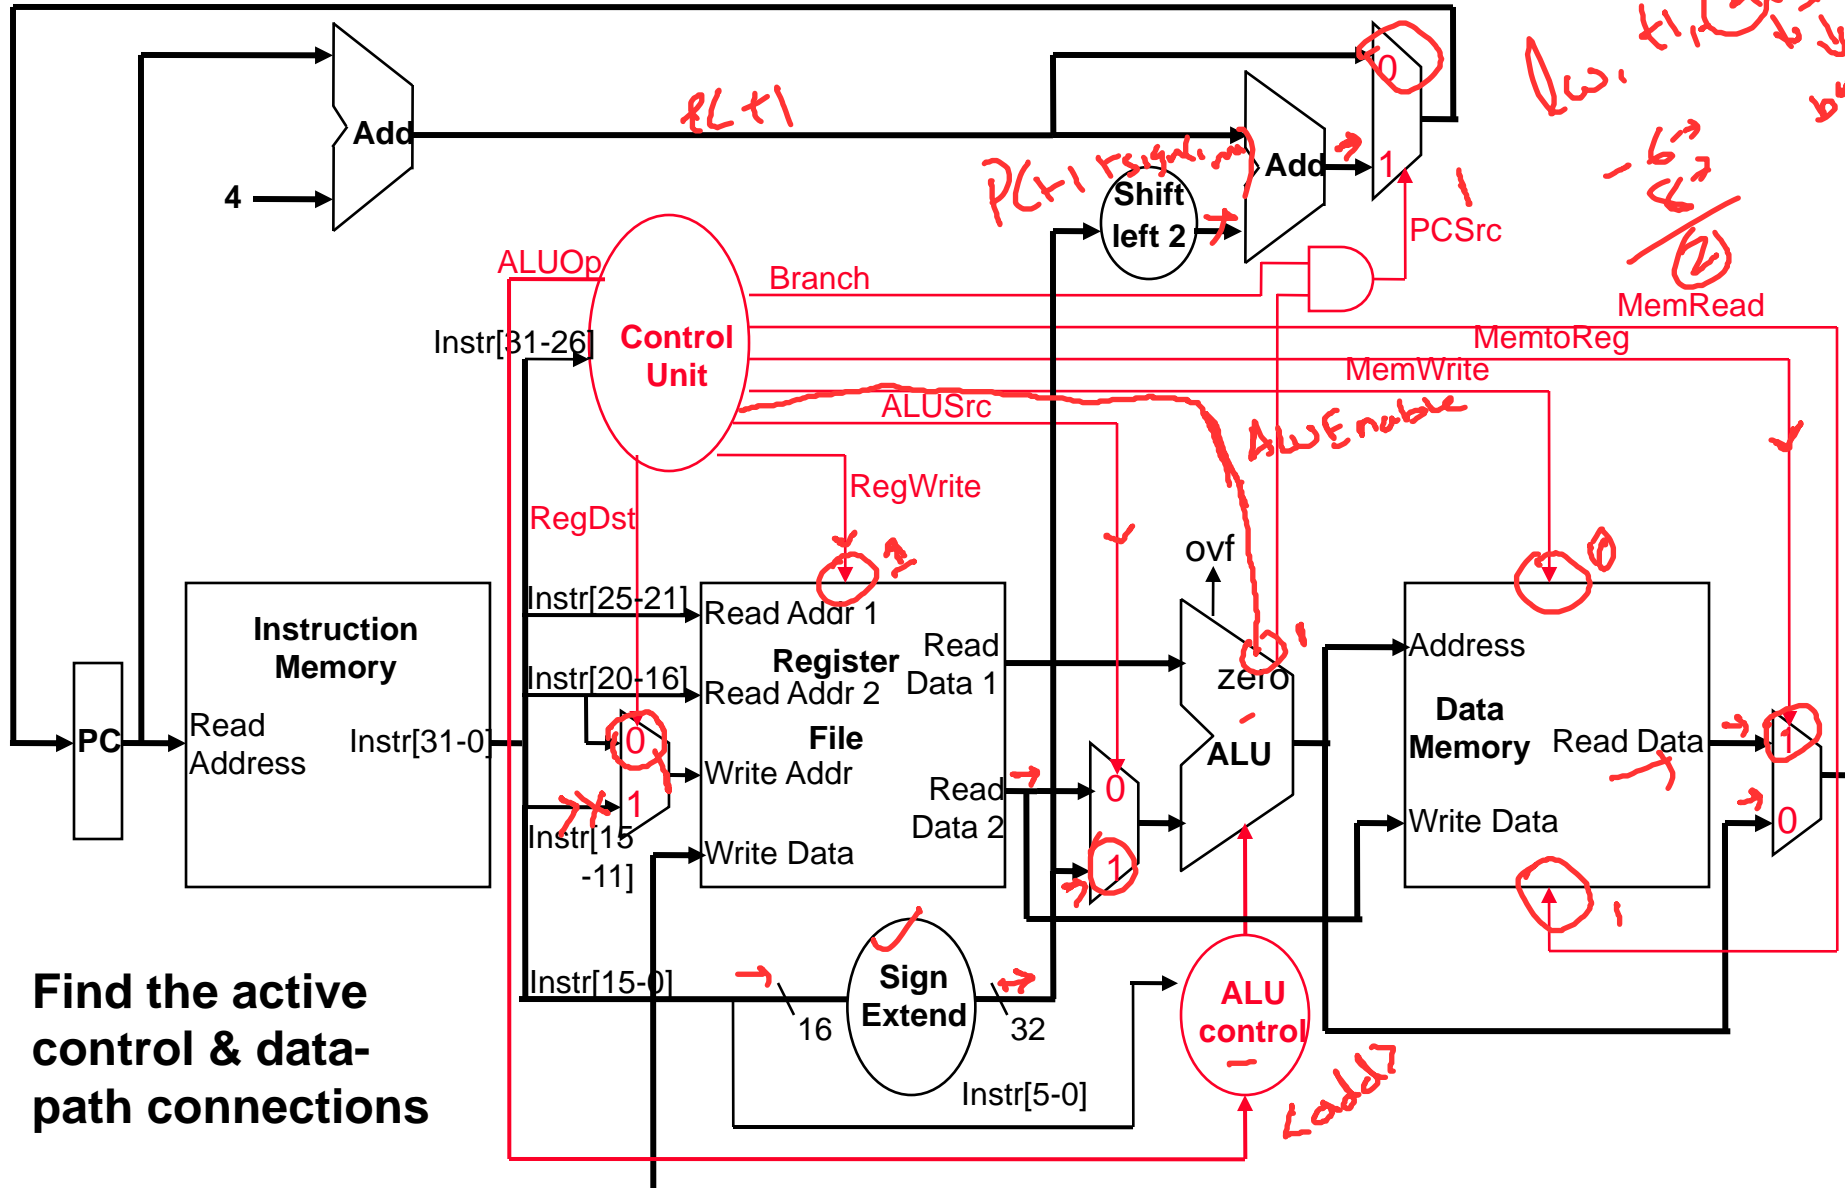
dp id ss st
/ok ss/rational



→ R-type Instruction Data/Control Flow

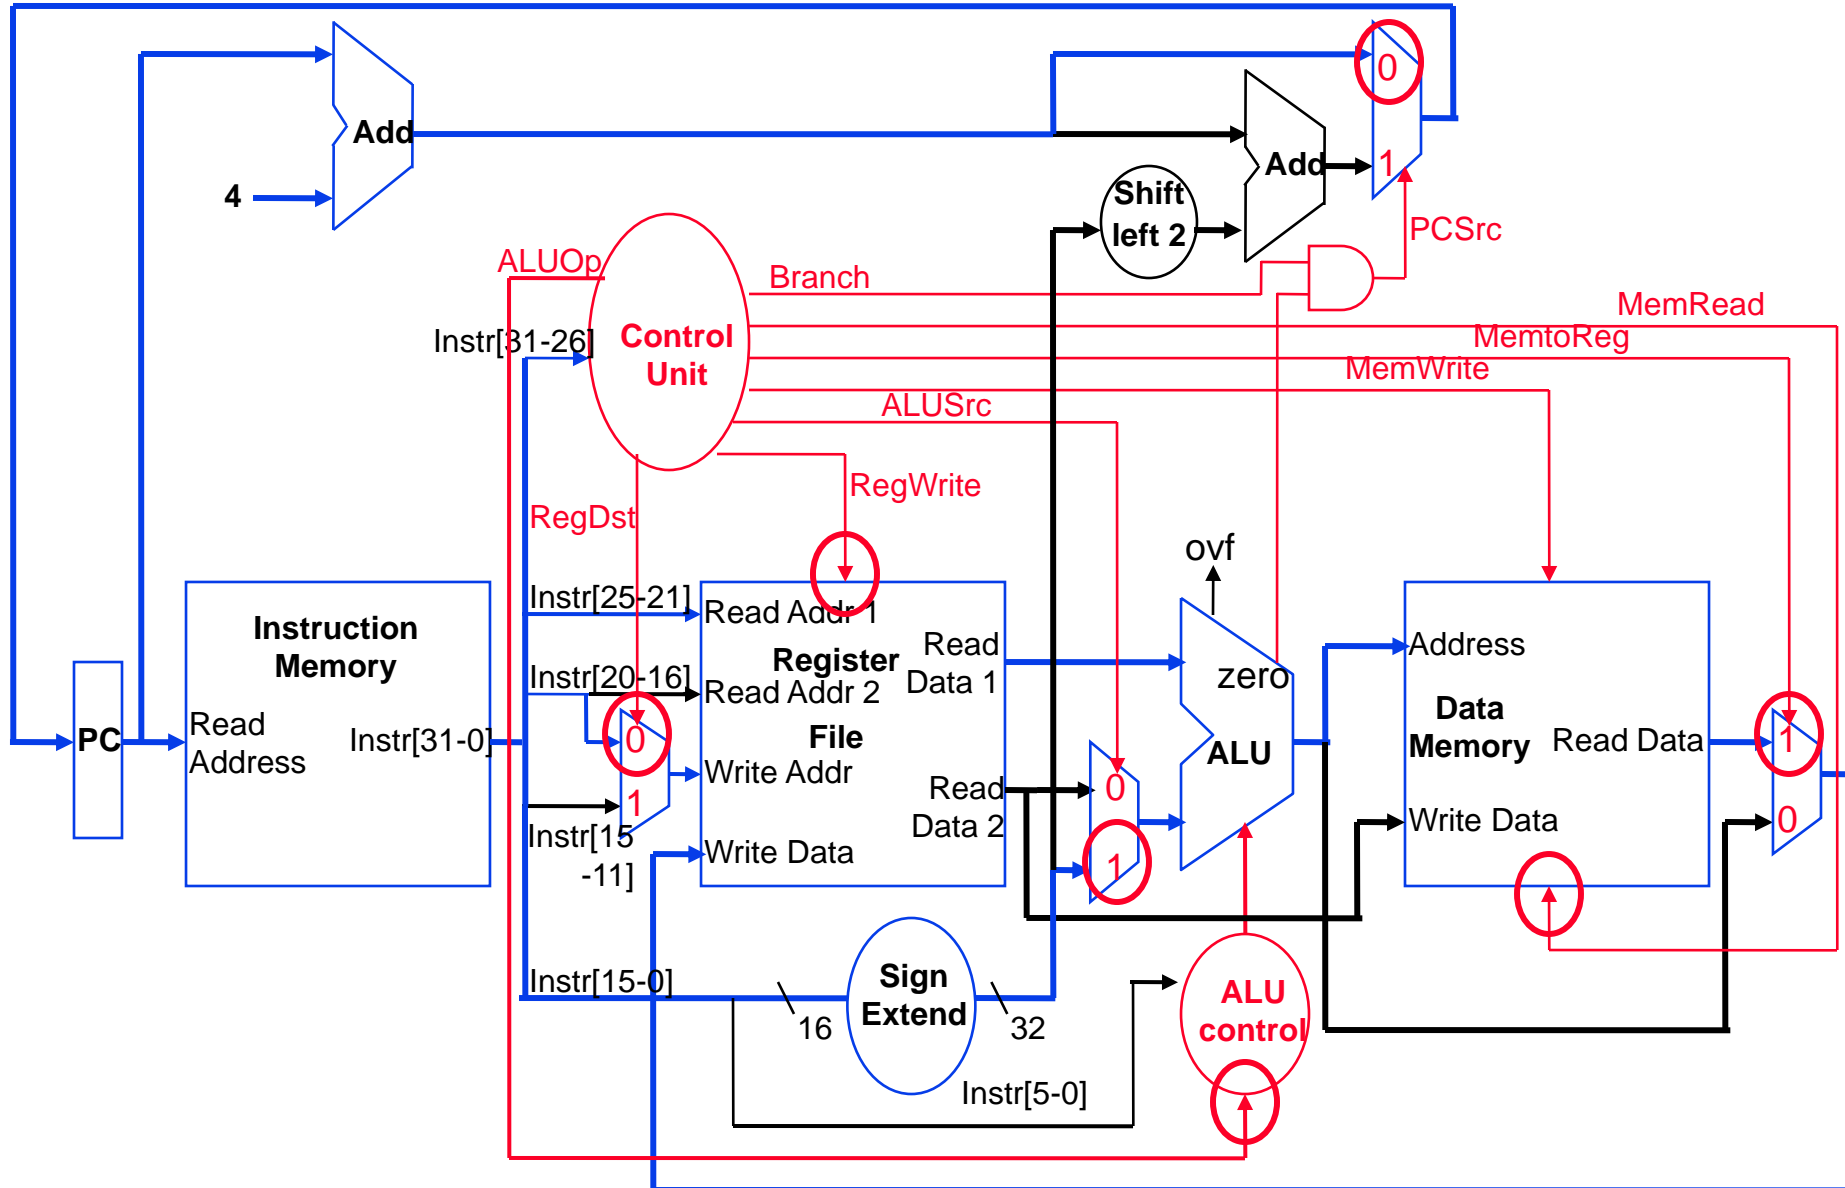


Load Word Instruction Data/Control Flow



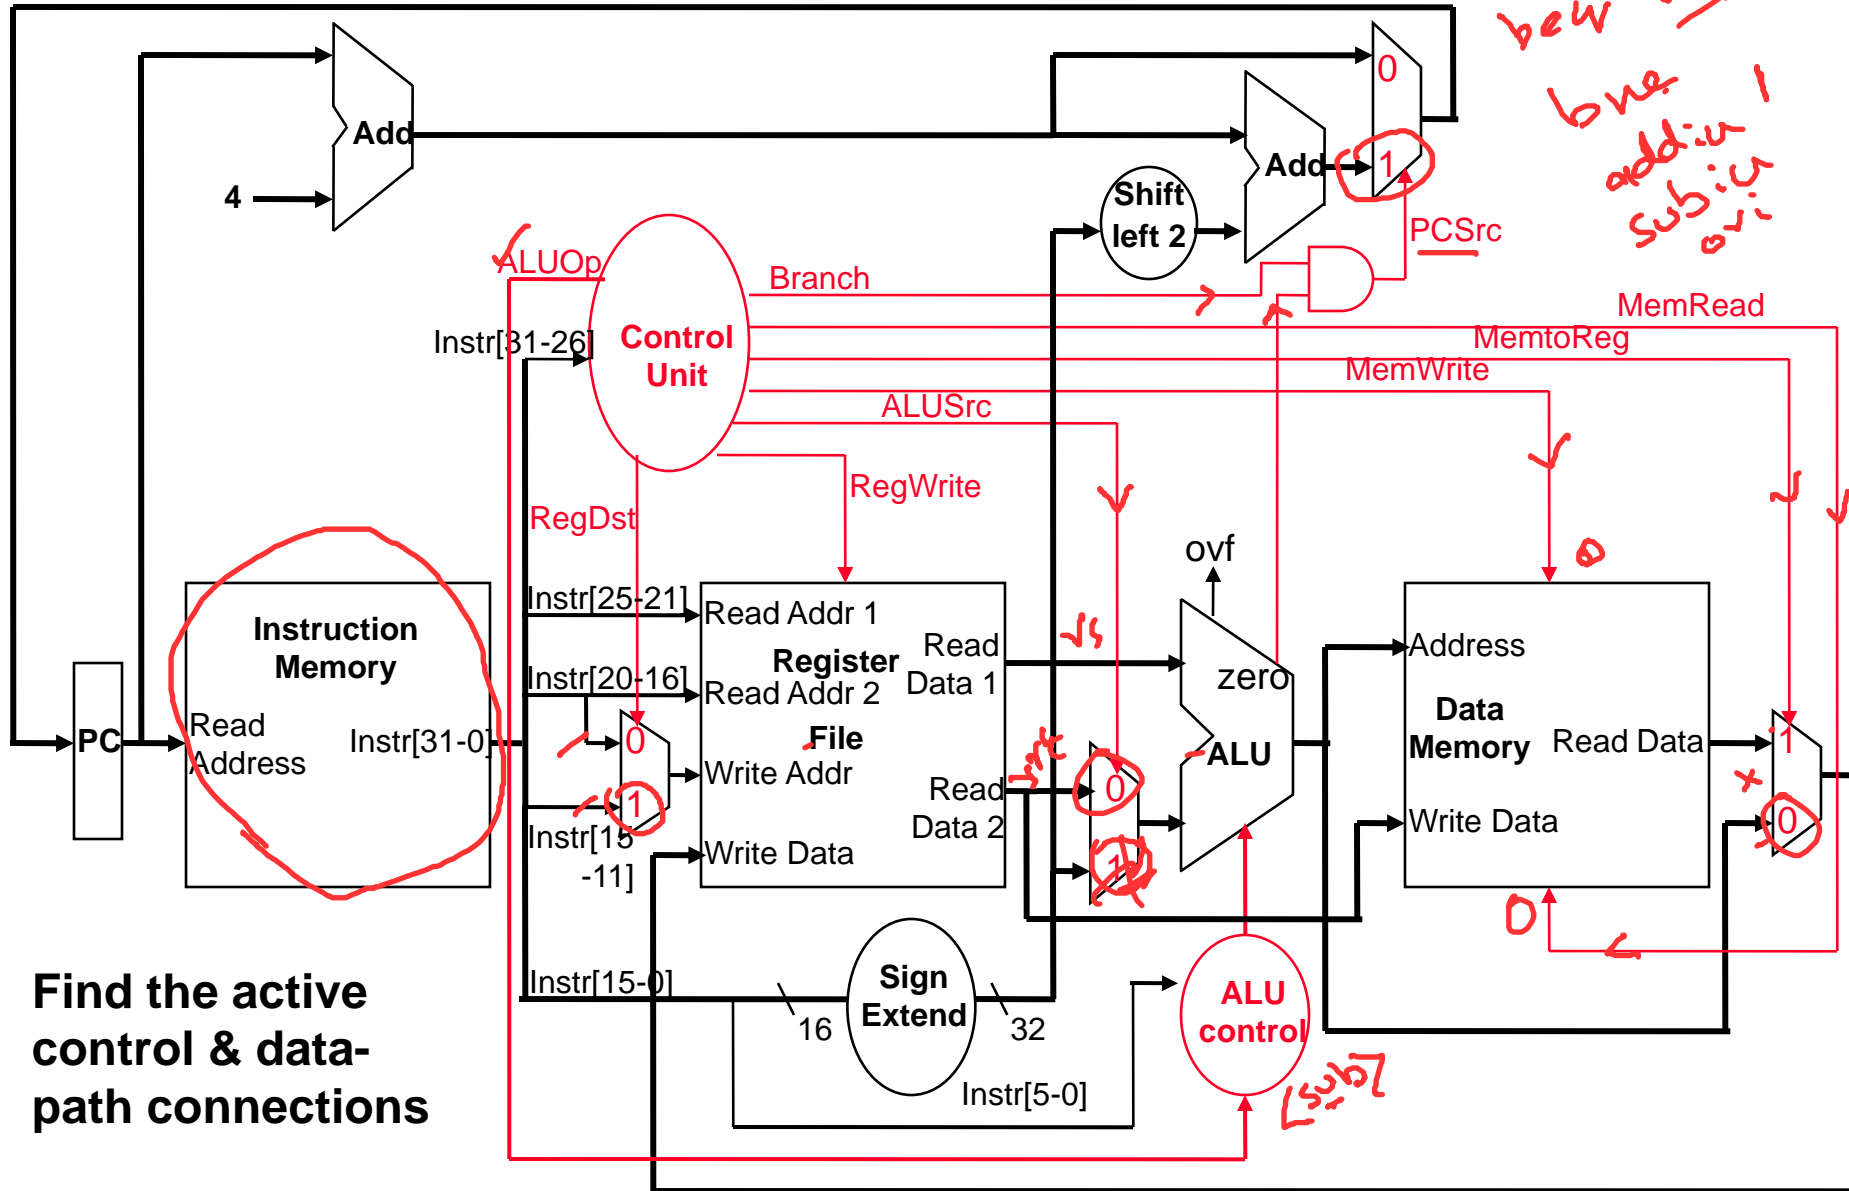
Find the active control & data-path connections

Load Word Instruction Data/Control Flow

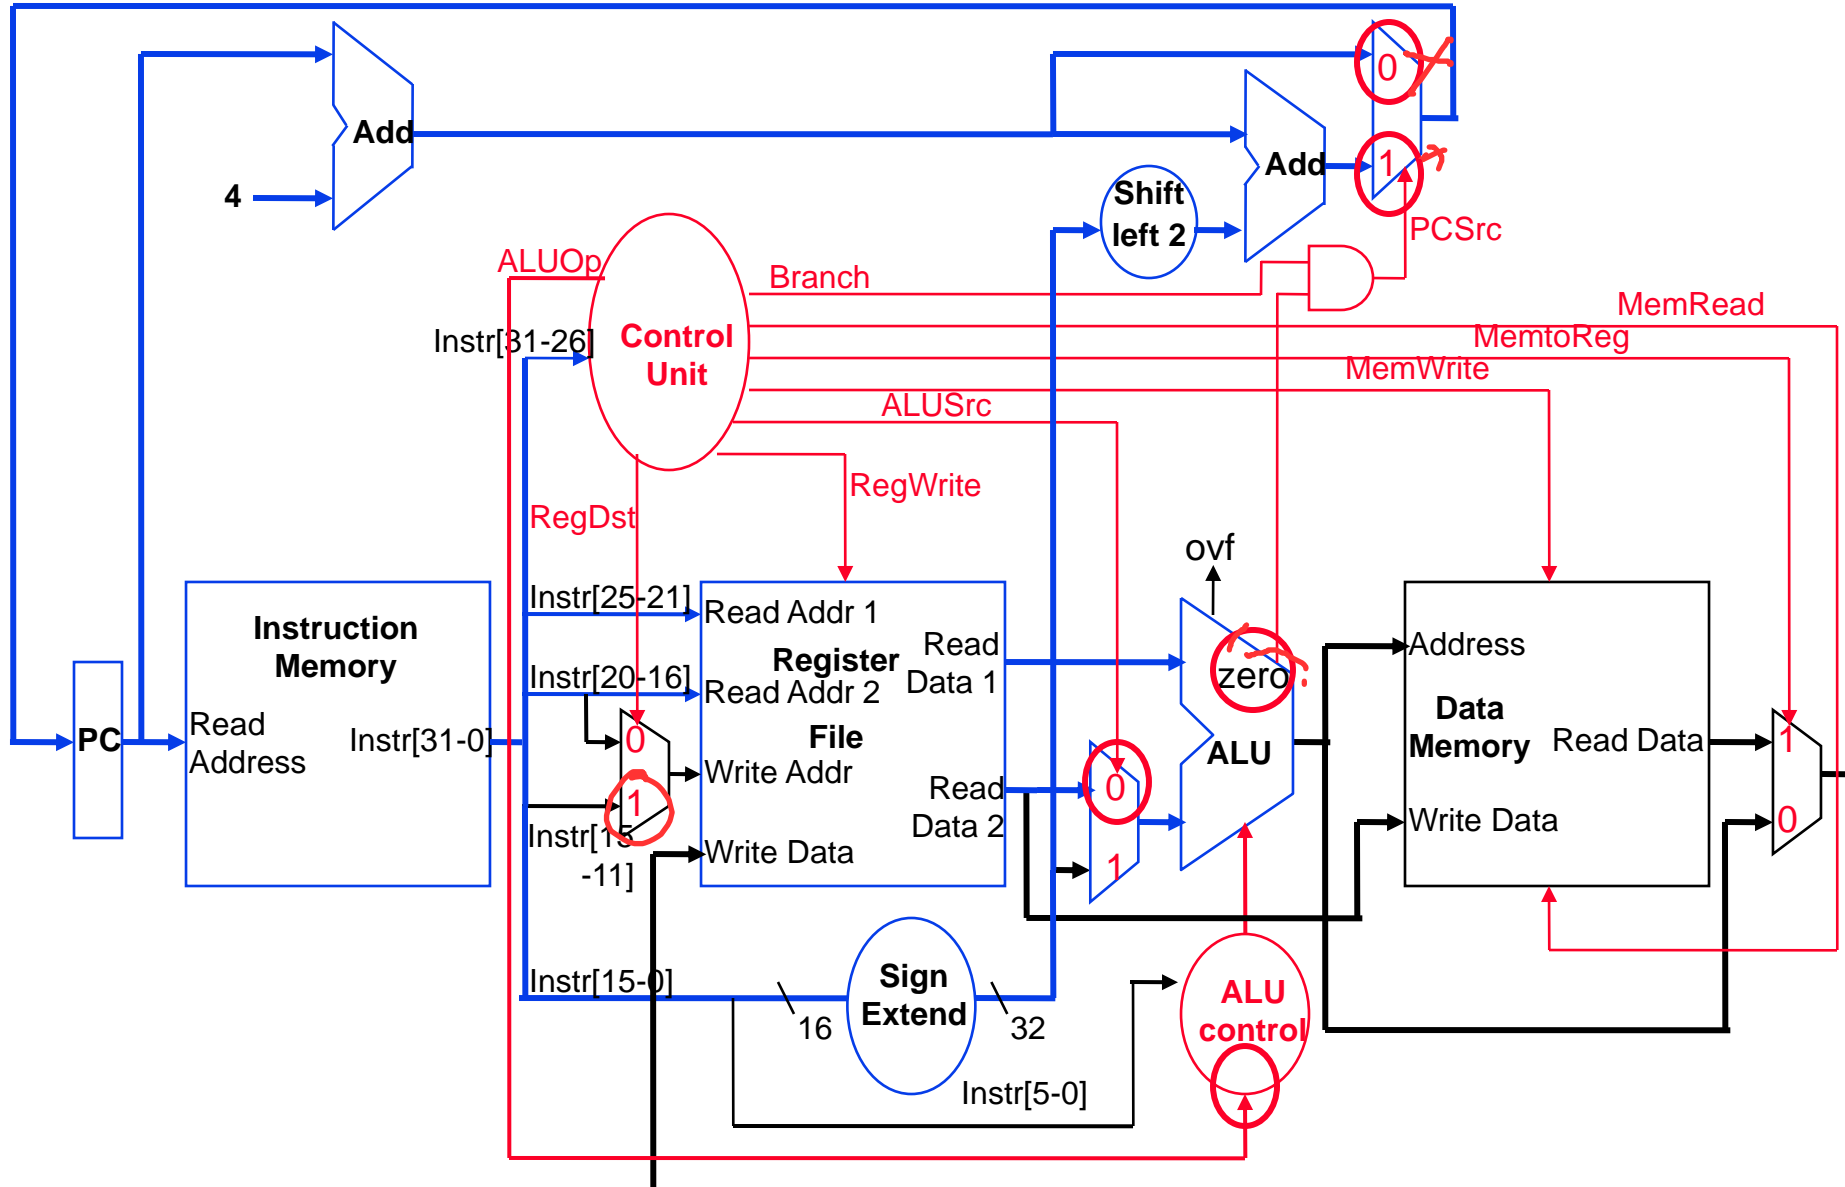




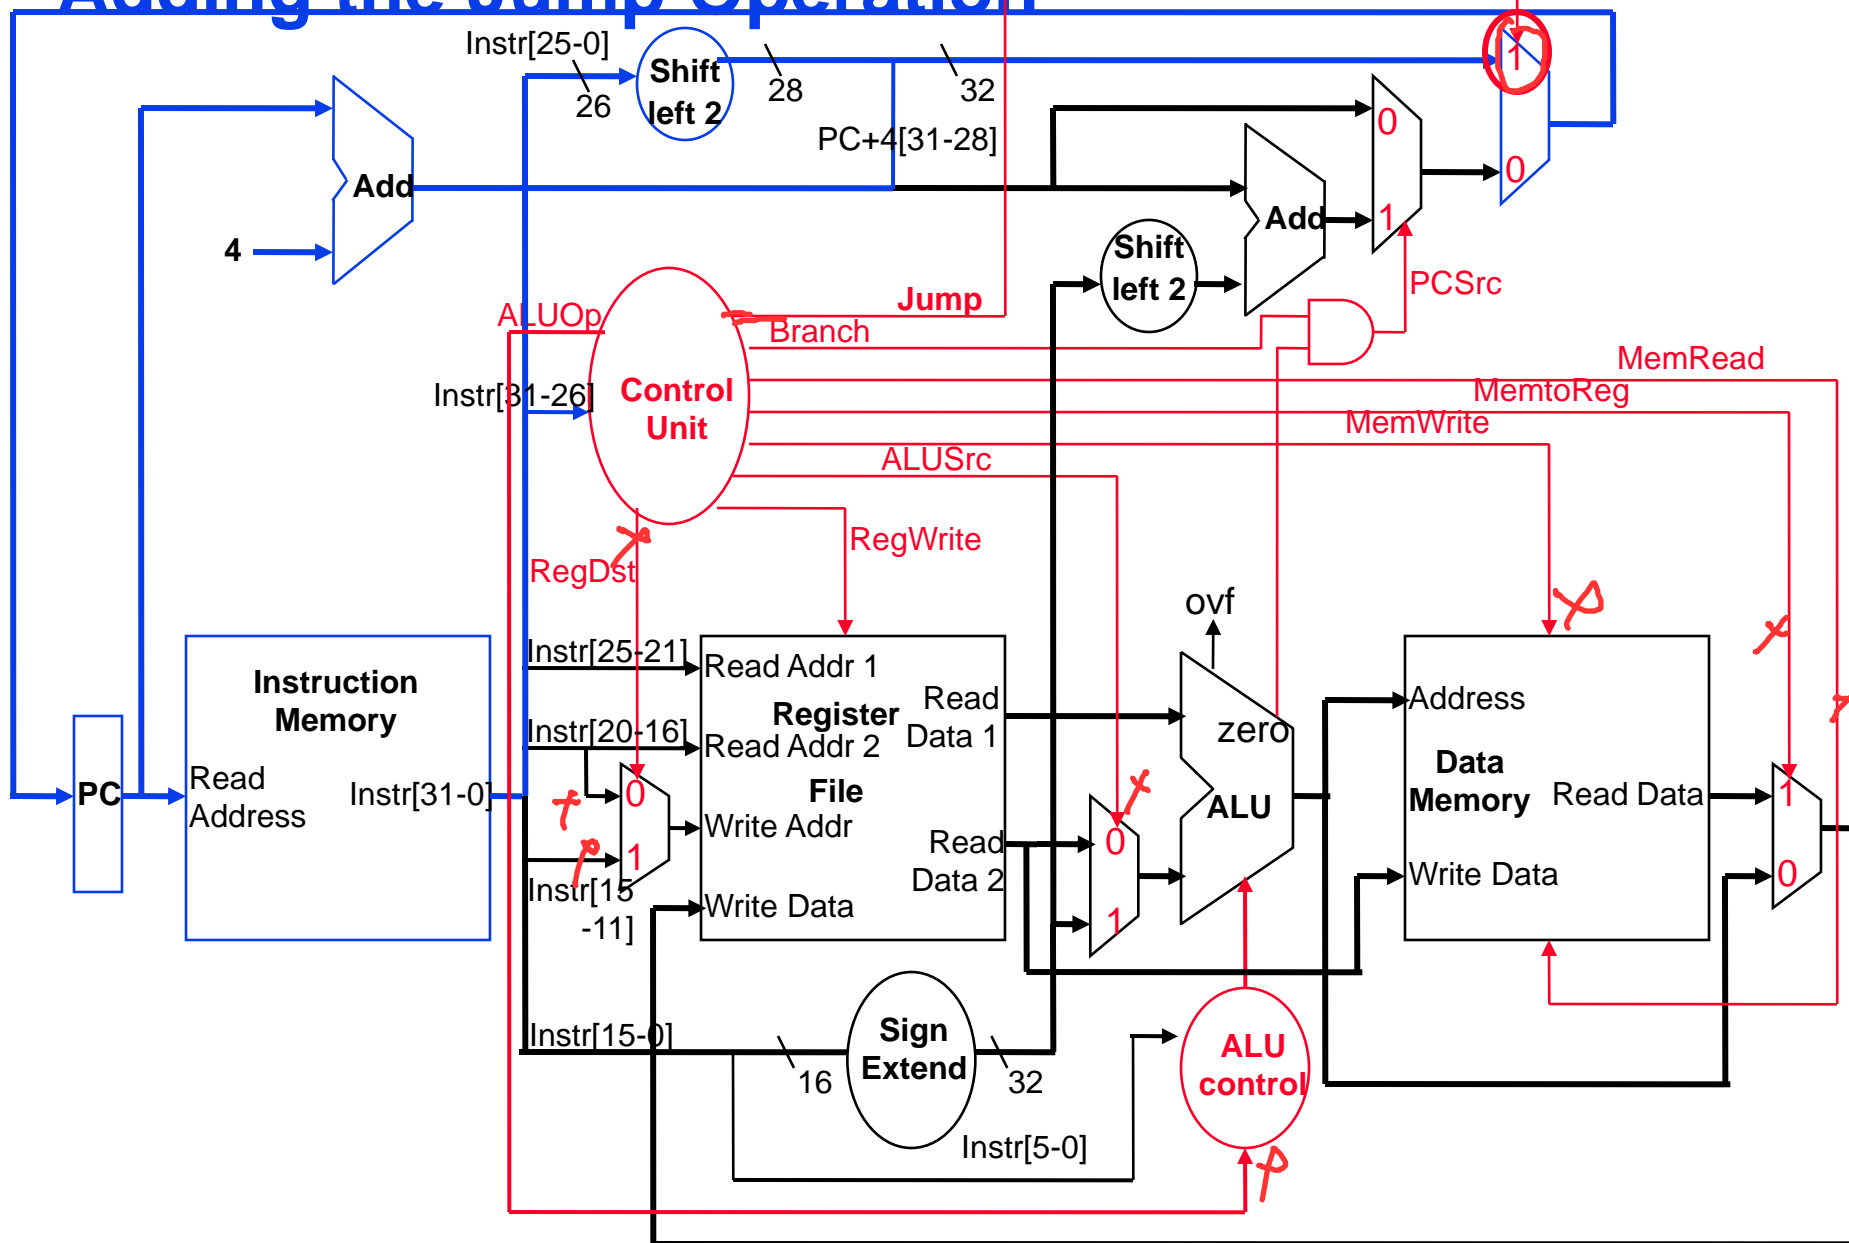
Branch Instruction Data/Control Flow



Branch Instruction Data/Control Flow



Adding the Jump Operation



Instruction Times (Critical Paths)

❑ What is the clock cycle time (assuming negligible delays for muxes, control unit, sign extend, PC access, shift left 2, wires, setup and hold times) but with:

- Instruction and Data Memory (200 ps)
- ALU and adders (200 ps)
- Register File access (reads or writes) (100 ps)

Instr.	I Mem	Reg Rd	ALU Op	D Mem	Reg Wr	Total
R-type	200	100	200	—	100	600 ps
load	200	100	200	200	100	800 ps
store	200	100	200	200	—	700 ps
beq	200	100	200	0	0	500 ps
jump	200	0	0	0	0	200 ps

800×10^{-9}
 8×10^{-12}
 25×10^{-12}
 25×10^{-12}

Instruction Critical Paths


❑ What is the clock cycle time (assuming negligible delays for muxes, control unit, sign extend, PC access, shift left 2, wires, setup and hold times) but with:

- Instruction and Data Memory (200 ps)
- ALU and adders (200 ps)
- Register File access (reads or writes) (100 ps)

Instr.	I Mem	Reg Rd	ALU Op	D Mem	Reg Wr	Total
R-type	200	100	200		100	600
→ load	200	100	200	200	100	800
store	200	100	200	200		700
beq	200	100	200			500
jump	200					200

Single Cycle Disadvantages & Advantages

- ❑ Uses the clock cycle inefficiently – the clock cycle must be timed to accommodate the **slowest** instruction
 - especially problematic for more complex instructions like floating point multiply



- ❑ May be wasteful of area since some functional units (e.g., adders) must be duplicated since they can not be shared during a clock cycle


but

- ❑ Is simple and easy to understand

How Can We Make It Faster?

- ❑ Start fetching and executing the next instruction before the current one has completed
 - **Pipelining** – all modern processors are pipelined for performance
 - Remember *the* performance equation:
$$\text{CPU time} = \text{CPI} * \text{CC} * \text{IC}$$
- ❑ Under *ideal* conditions and with a large number of instructions, the speedup from pipelining is approximately equal to the number of pipe stages
 - A five stage pipeline is nearly five times faster because the CC is nearly five times faster
- ❑ Fetch (and execute) more than one instruction at a time
 - Superscalar processing – stay tuned

The Five Stages of Load Instruction

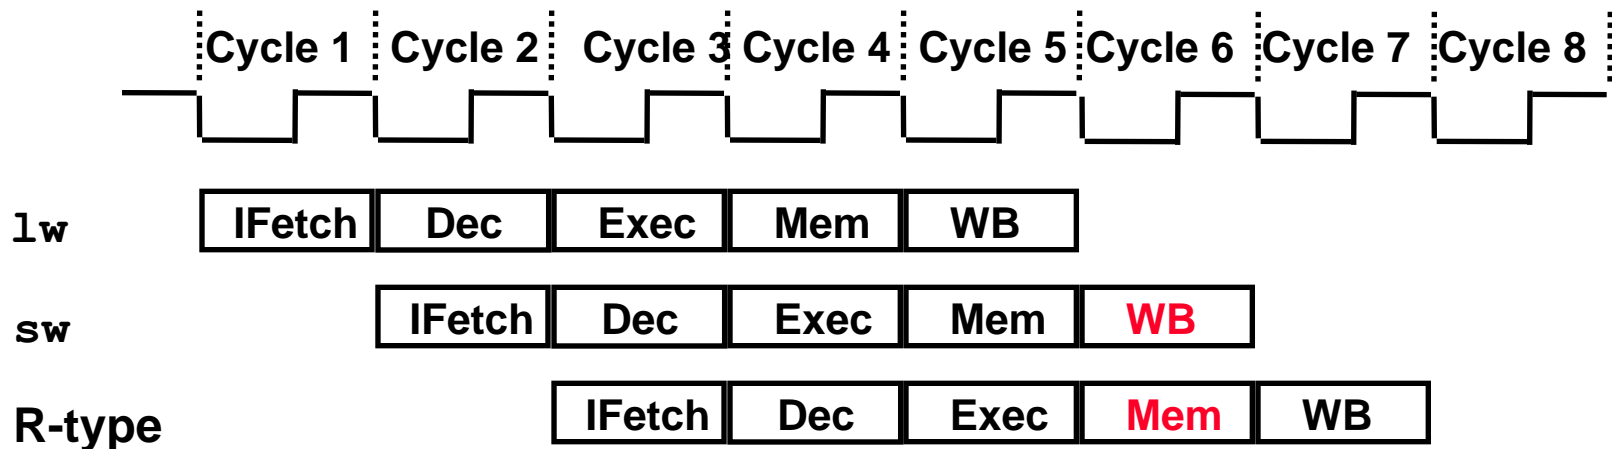


- ❑ IFetch: Instruction Fetch and Update PC
- ❑ Dec: Registers Fetch and Instruction Decode
- ❑ Exec: Execute R-type; calculate memory address
- ❑ Mem: Read/write the data from/to the Data Memory
- ❑ WB: Write the result data into the register file

A Pipelined RISC Processor

❑ Start the **next** instruction before the current one has completed


- improves **throughput** - total amount of work done in a given time
- instruction **latency** (execution time, delay time, response time - time from the start of an instruction to its completion) is *not* reduced



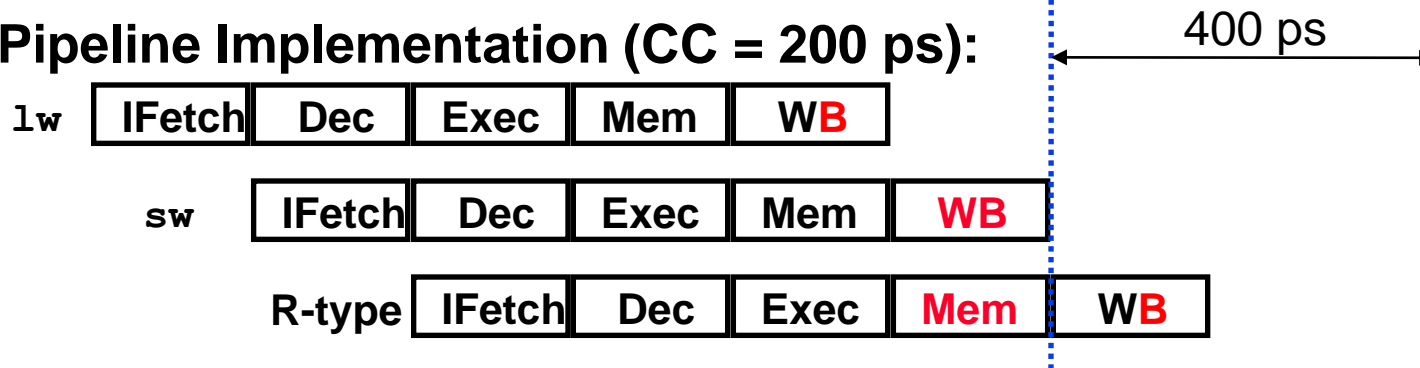
- clock cycle (pipeline stage time) is limited by the **slowest** stage
 - for some stages don't need the whole clock cycle (e.g., WB)
 - for some instructions, some stages are **wasted** cycles (i.e., nothing is done during that cycle for that instruction)

Single Cycle versus Pipeline

Single Cycle Implementation (CC = 800 ps):



Pipeline Implementation (CC = 200 ps):



- ❑ To complete an entire instruction in the pipelined case takes 1000 ps (as compared to 800 ps for the single cycle case). Why ?
- ❑ How long does each take to complete 1,000,000 adds ?

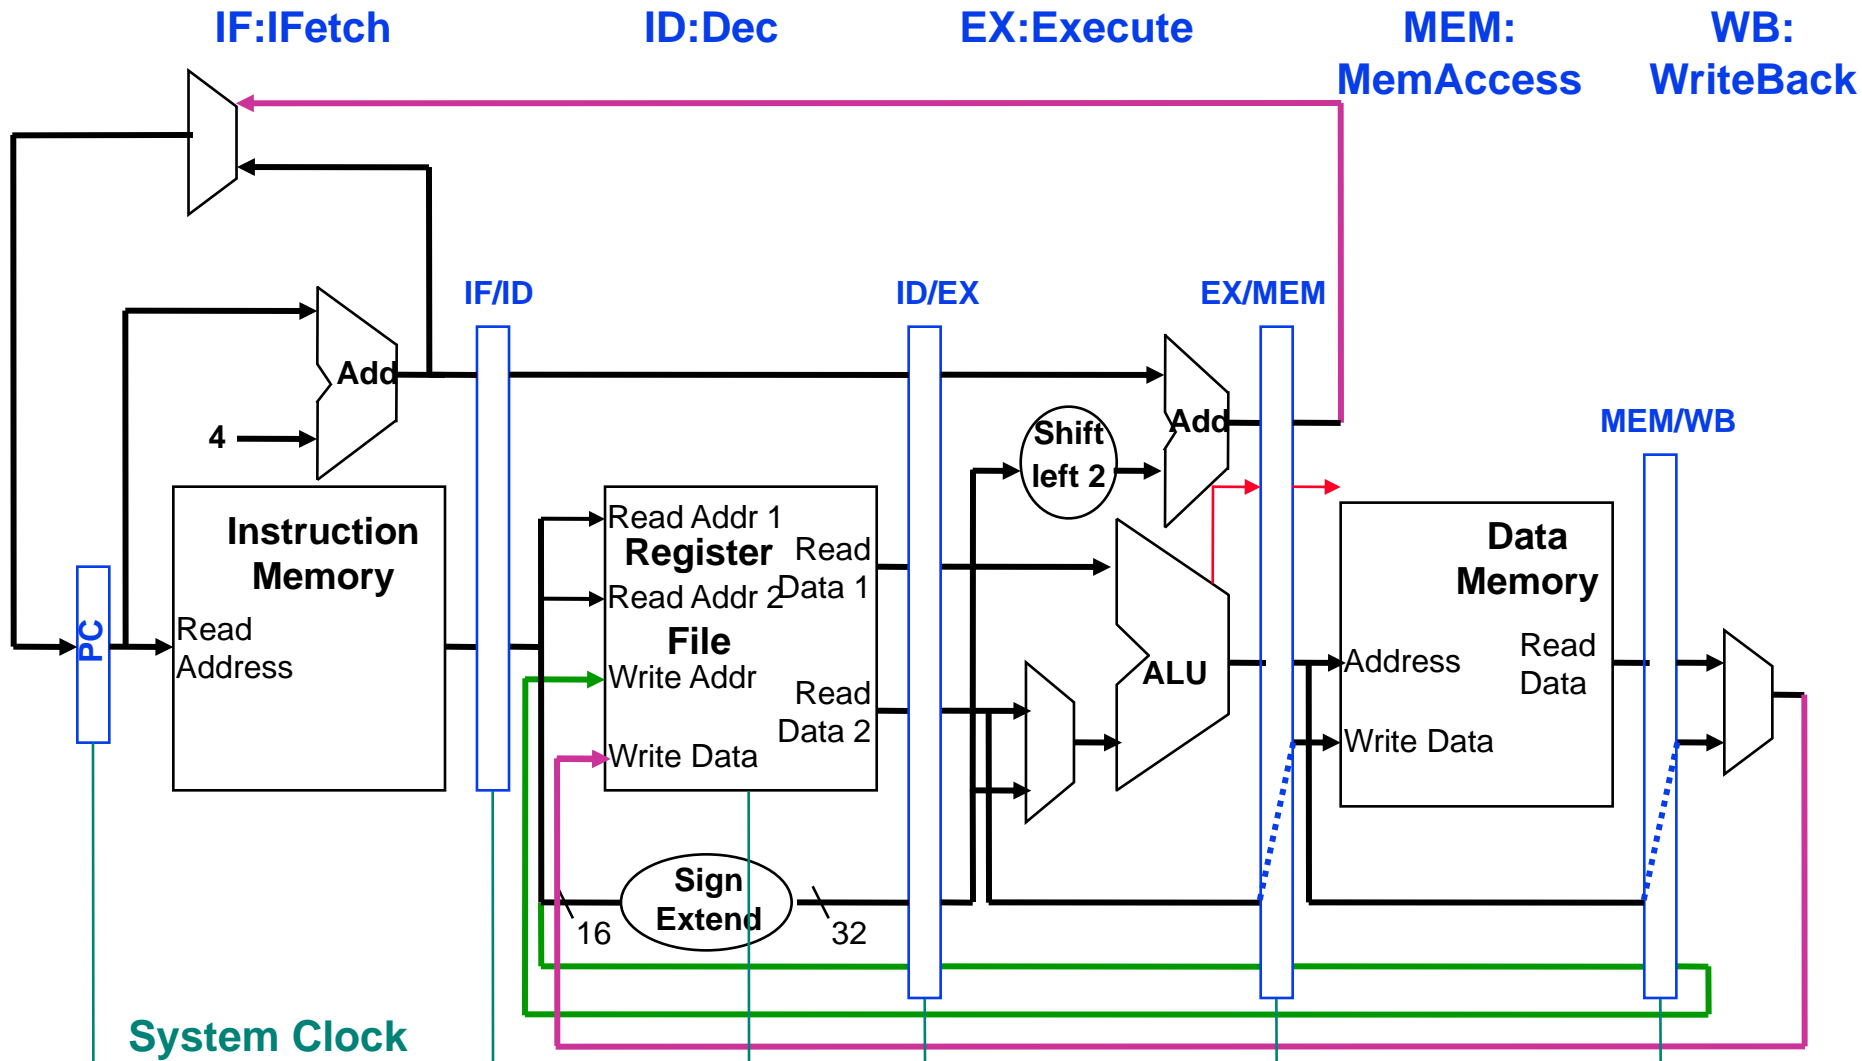
Pipelining the RISC ISA

□ What makes it easy

- all instructions are the same length (32 bits)
 - can fetch in the 1st stage and decode in the 2nd stage
- few instruction formats (only 3) with **symmetry** across formats
 - can begin reading register file in 2nd stage
- memory operations occur only in loads and stores
 - can use the execute stage to calculate memory addresses
- each instruction writes at most one result (i.e., changes the machine state) and does it in the last few pipeline stages (MEM or WB)
- operands must be aligned in memory so a single data transfer takes only one data memory access

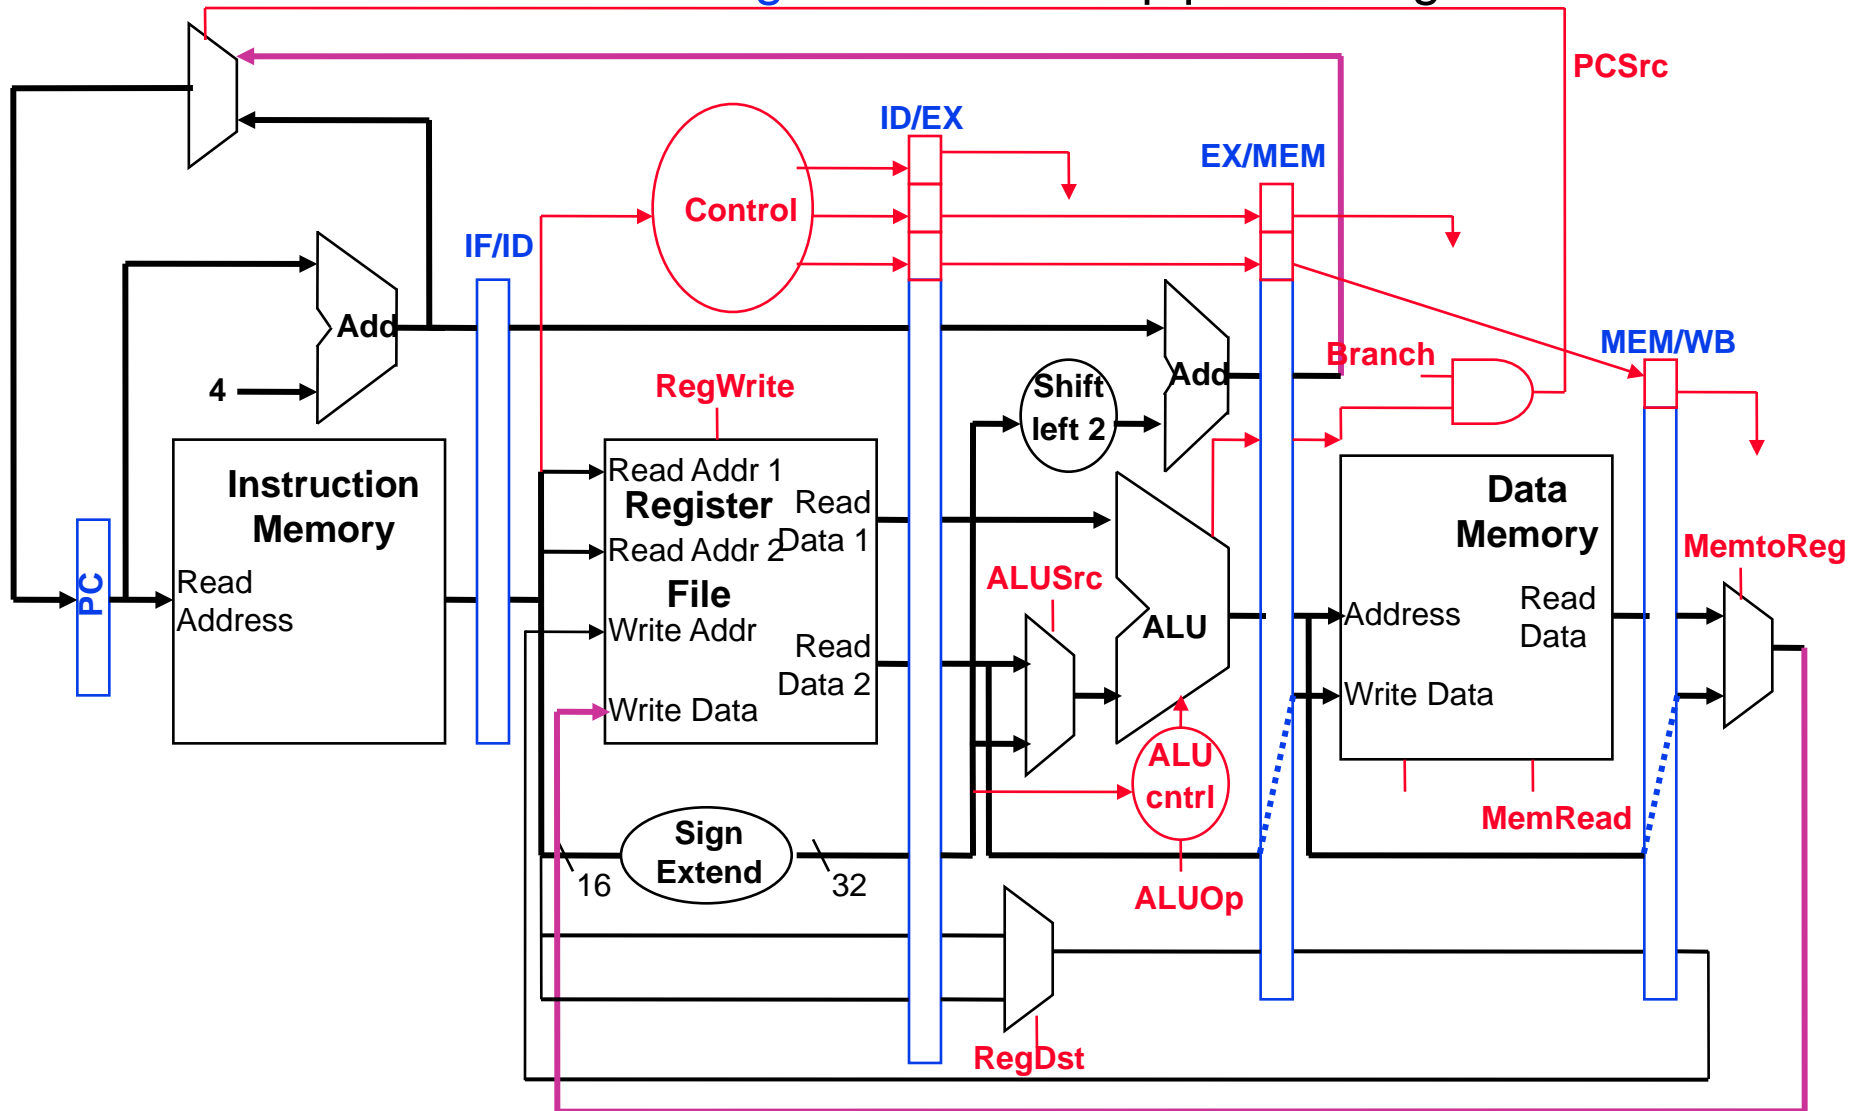
RISC Pipeline Datapath Additions/Mods

❑ State registers between each pipeline stage to **isolate** them



RISC Pipeline Control Path Modifications

- ❑ All control signals can be determined during Decode
- and held in the **state registers** between pipeline stages




Pipeline Control

- ❑ IF Stage: read Instr Memory (always asserted) and write PC (on System Clock)
- ❑ ID Stage: no optional control signals to set

	EX Stage				MEM Stage			WB Stage	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Brch	Mem Read	Mem Write	Reg Write	Mem toReg
R	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

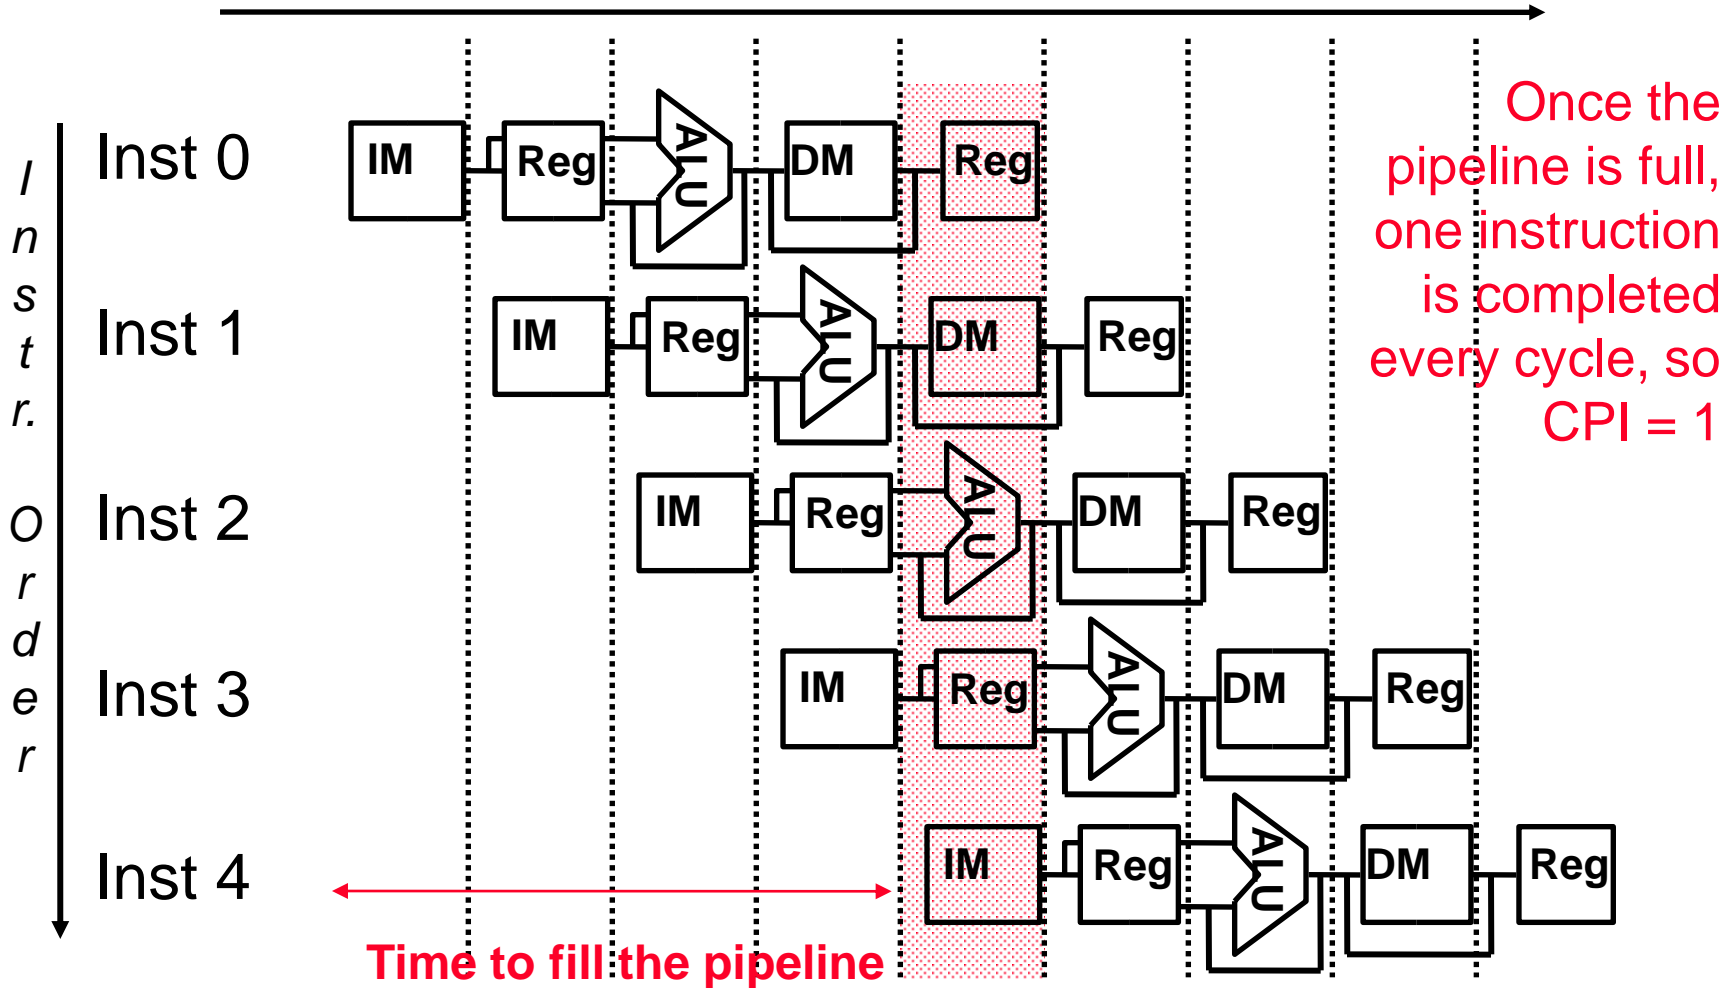
Graphically Representing RISC Pipeline



- Can help with answering questions like:
- How many cycles does it take to execute this code?
 - What is the ALU doing during cycle 4?
 - Is there a hazard, why does it occur, and how can it be fixed?

Why Pipeline? For Performance!

Time (clock cycles)



Can Pipelining Get Us Into Trouble?

❑ Yes: Pipeline Hazards

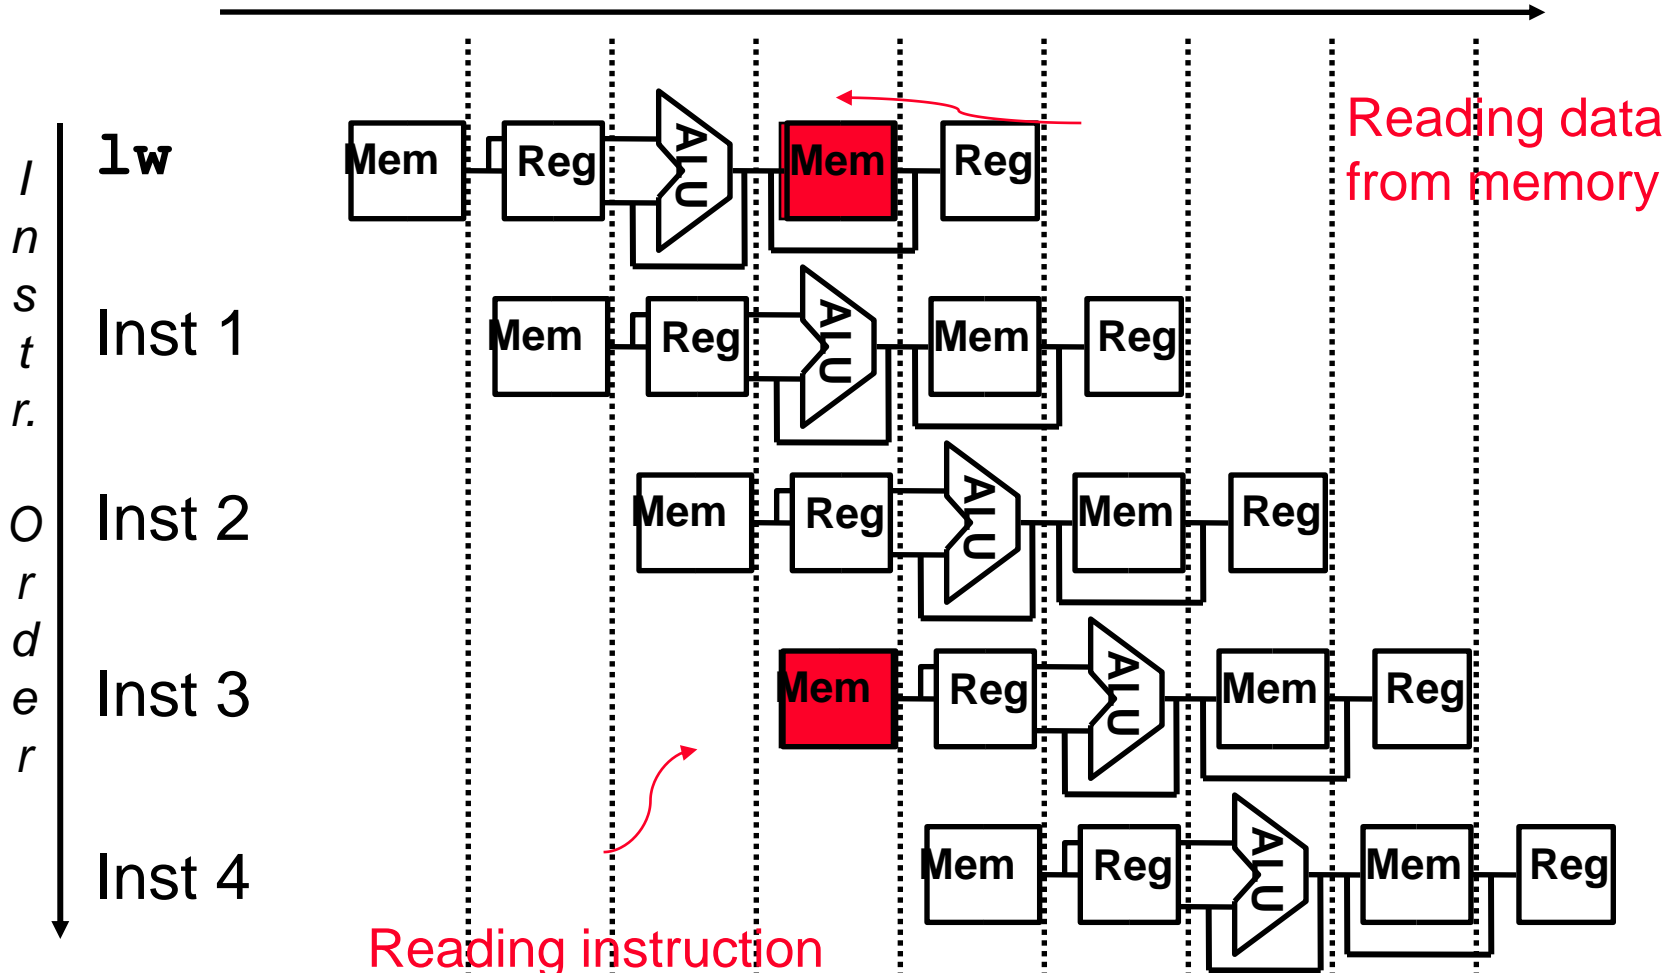
- **structural hazards**: attempt to use the same resource by two different instructions at the same time
- **data hazards**: attempt to use data before it is ready
 - An instruction's source operand(s) are produced by a prior instruction still in the pipeline
- **control hazards**: attempt to make a decision about program control flow before the condition has been evaluated and the new PC target address calculated
 - branch and jump instructions, exceptions

❑ Can usually resolve hazards by waiting

- pipeline control must **detect** the hazard
- and take action to **resolve** hazards

A Single Memory Would Be a Structural Hazard

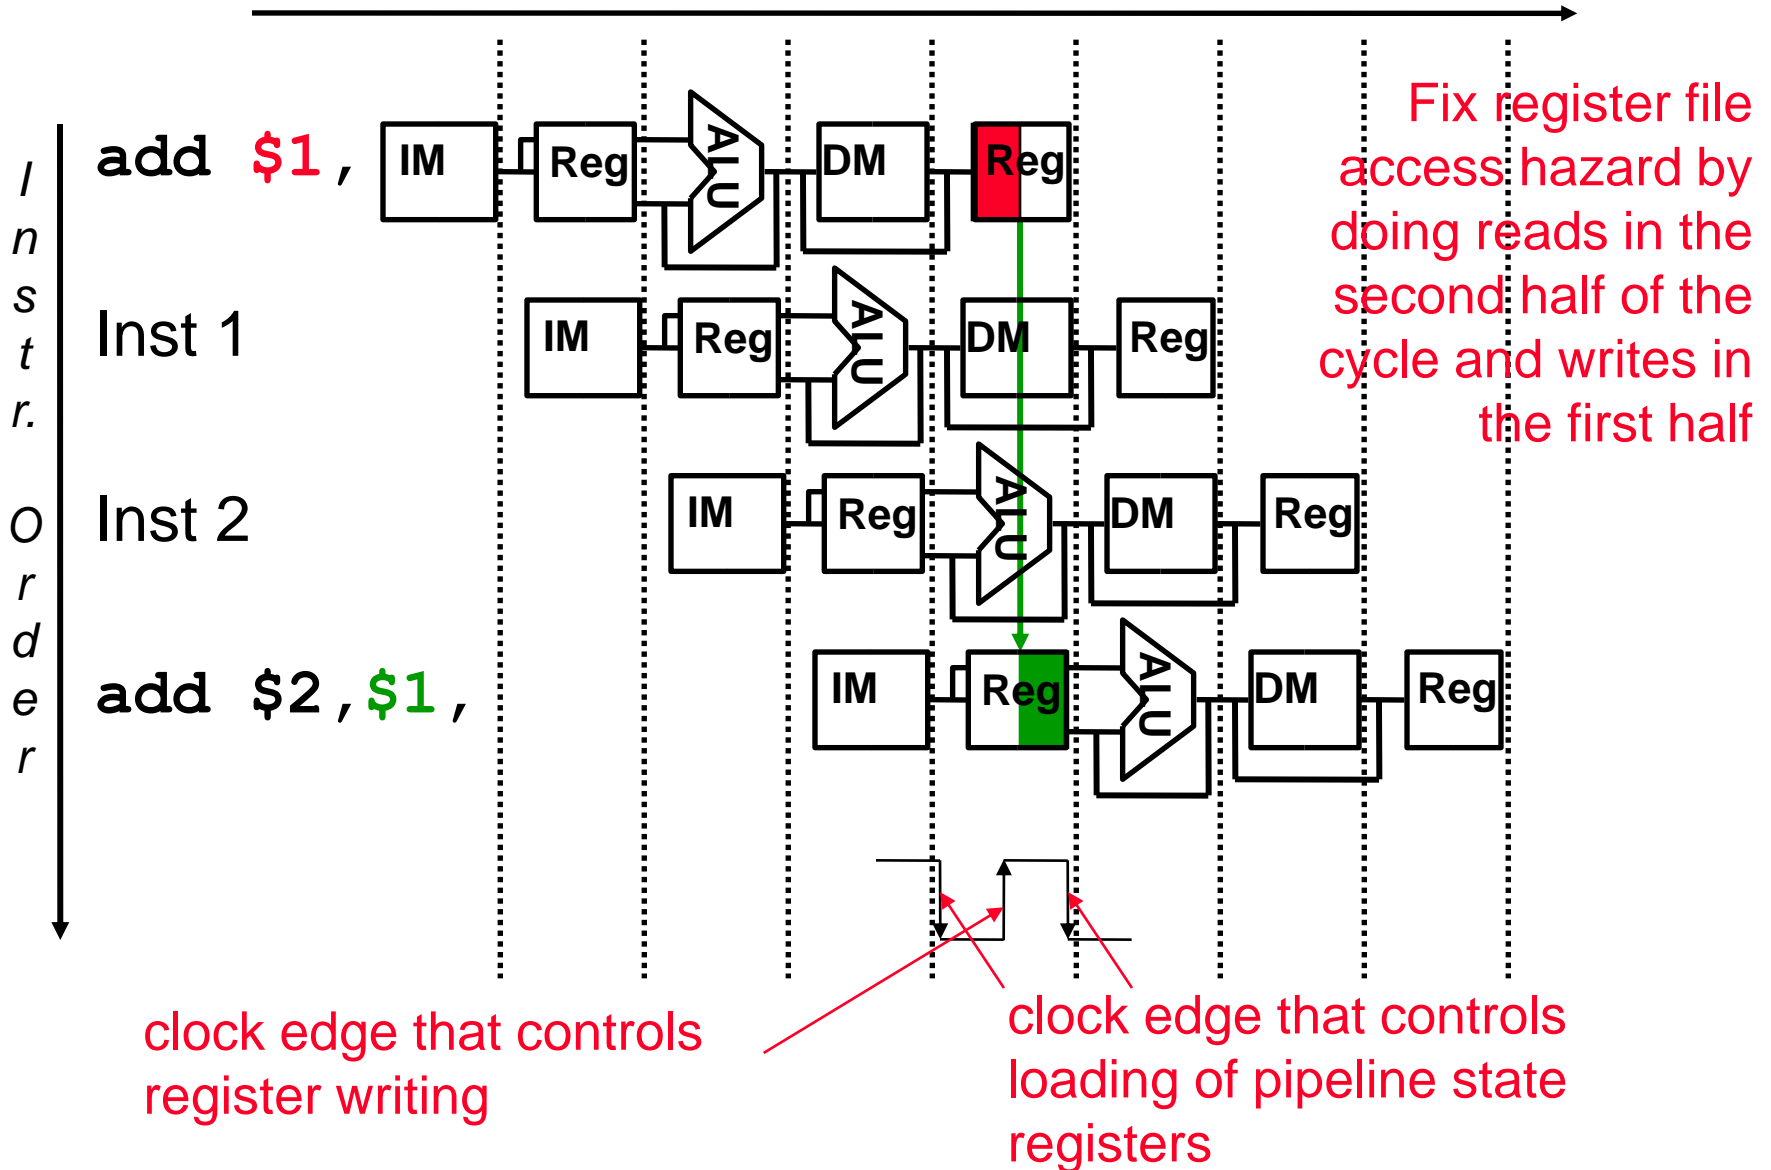
Time (clock cycles)



❑ Fix with separate instr and data memories (I\$ and D\$)

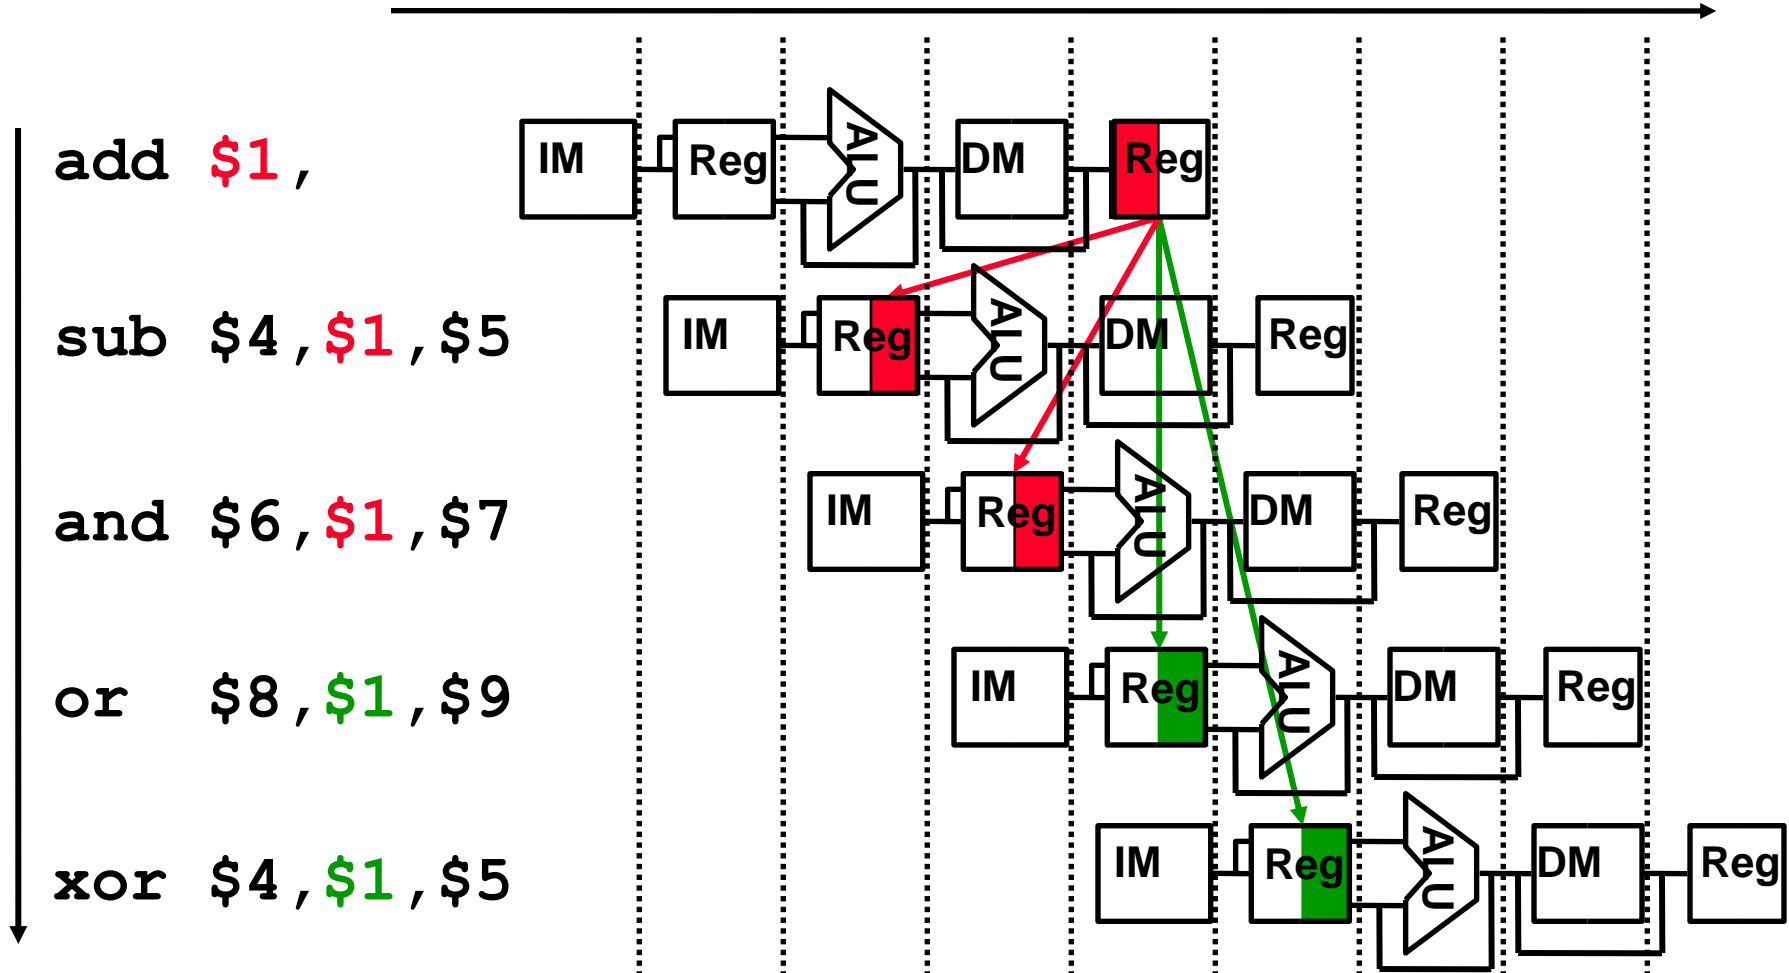
How About Register File Access?

Time (clock cycles)



Register Usage Can Cause Data Hazards

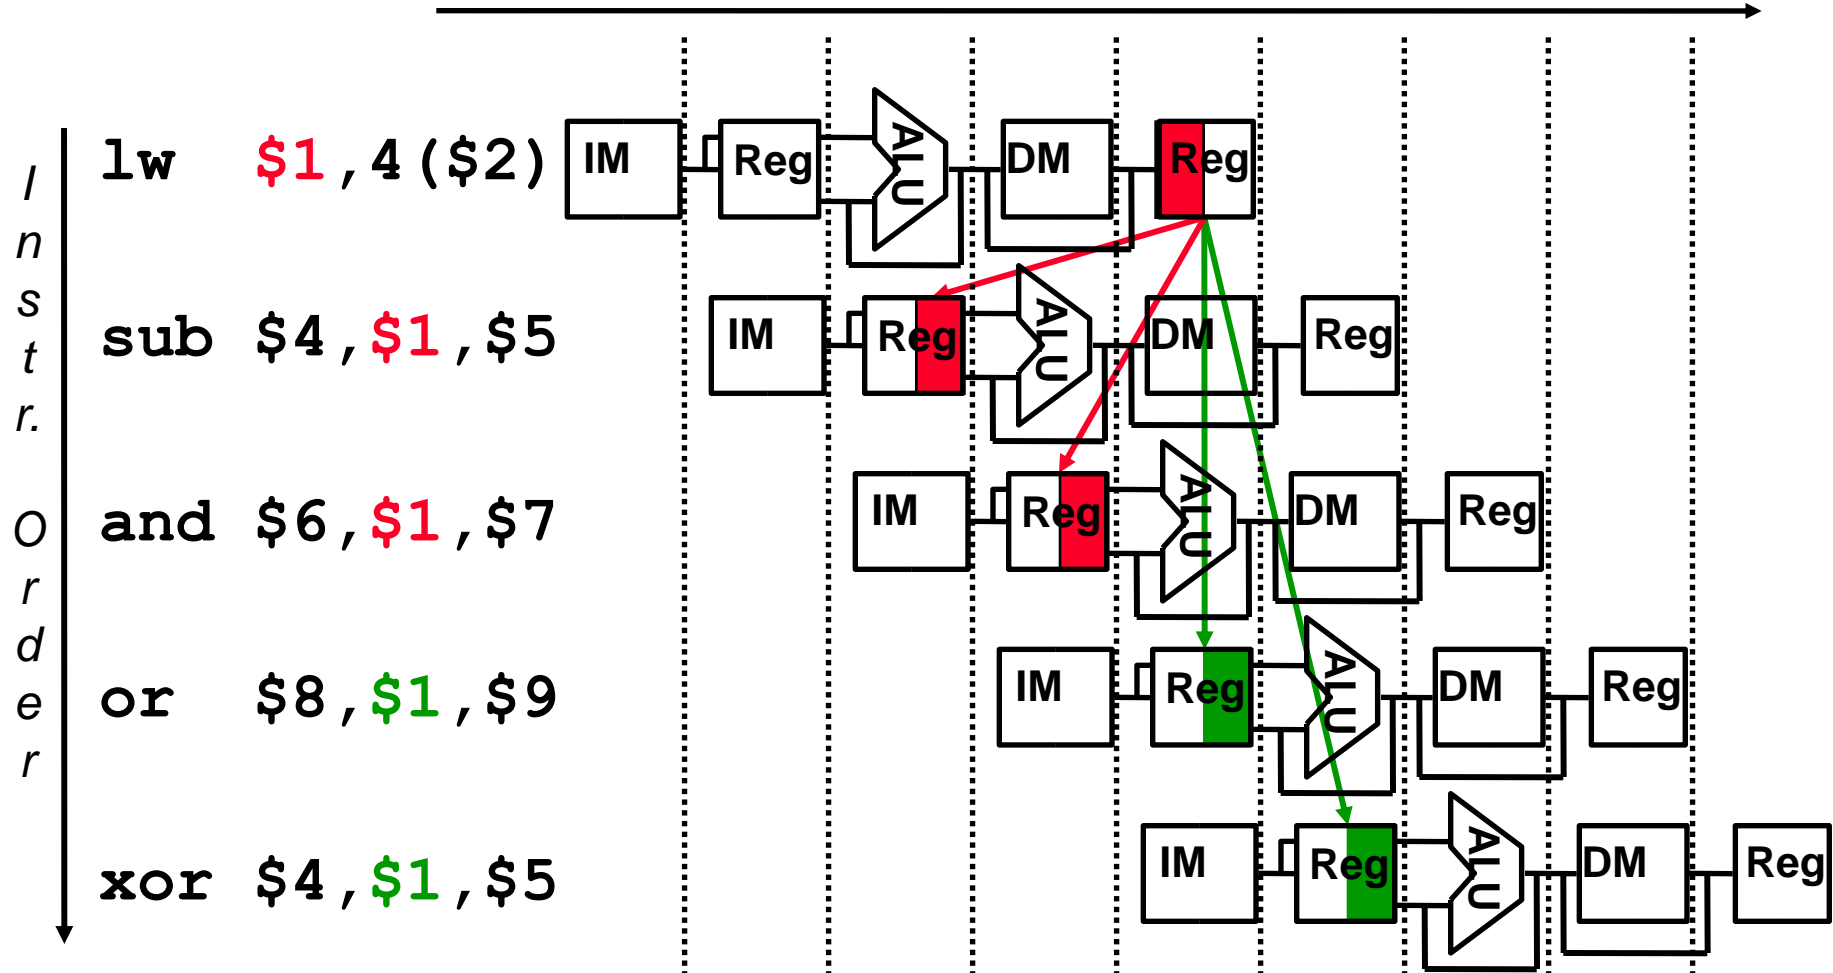
❑ Dependencies backward in time cause **hazards**



❑ Read after write data hazard (RAW in the code)

Loads Can Cause Data Hazards

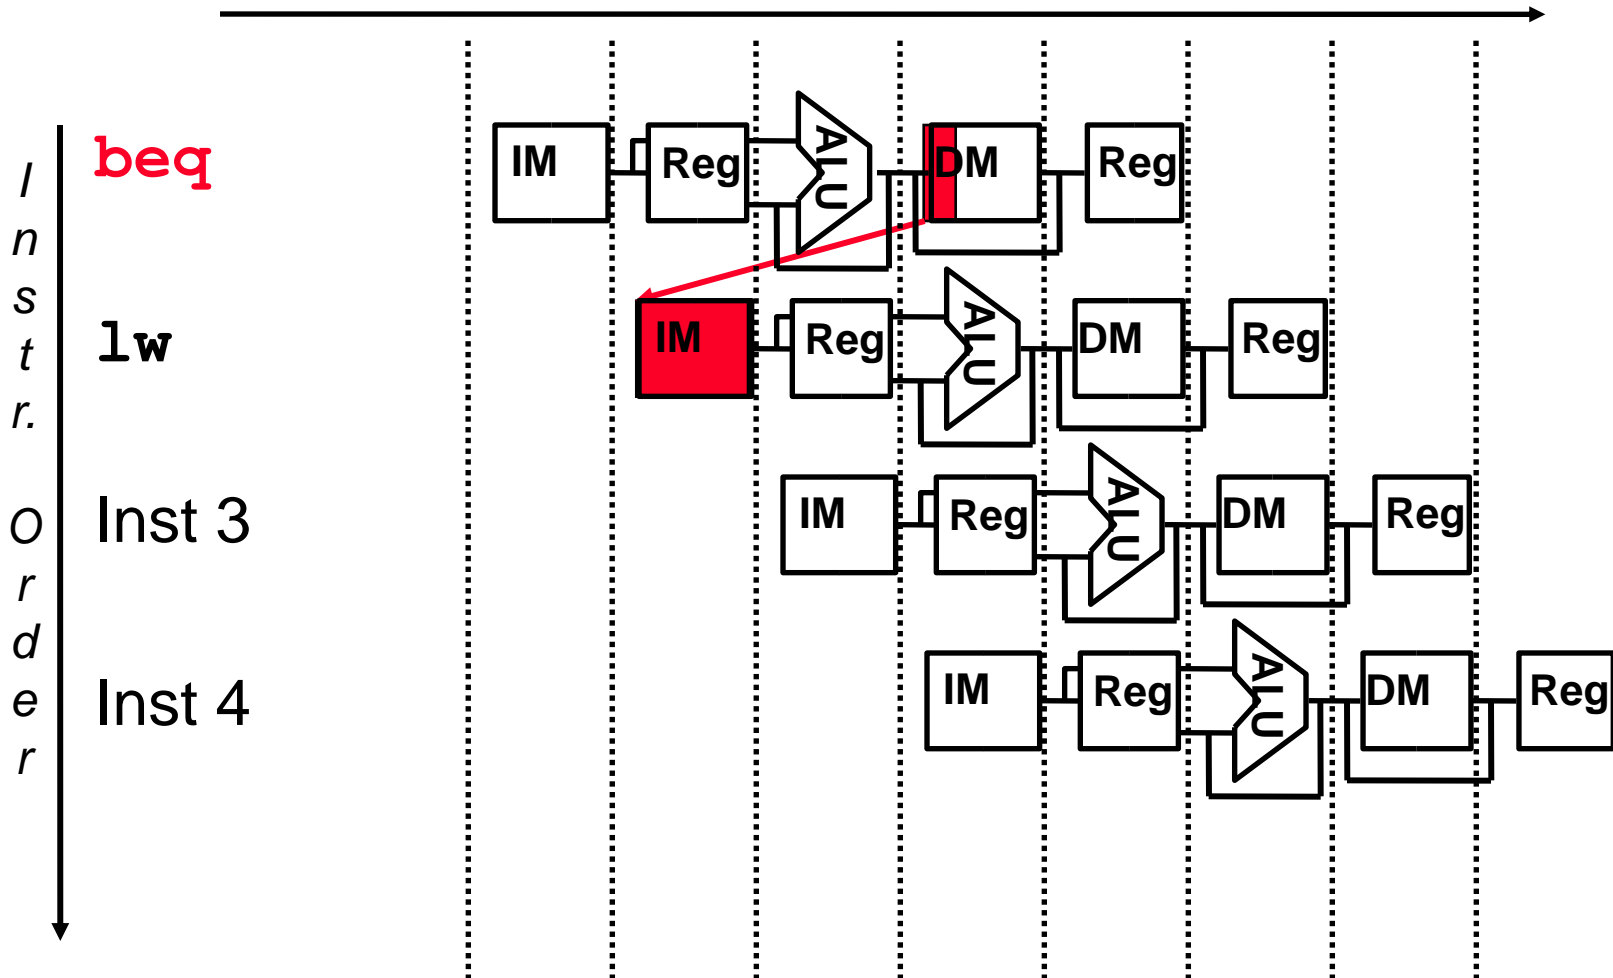
Dependencies backward in time cause hazards



Load-use data hazard

Branch Instructions Cause Control Hazards


❑ Dependencies backward in time cause **hazards**



Other Pipeline Structures Are Possible


❑ What about the (slow) multiply operation?

- Make the clock twice as slow or ...
- let it take two cycles (since it doesn't use the DM stage)



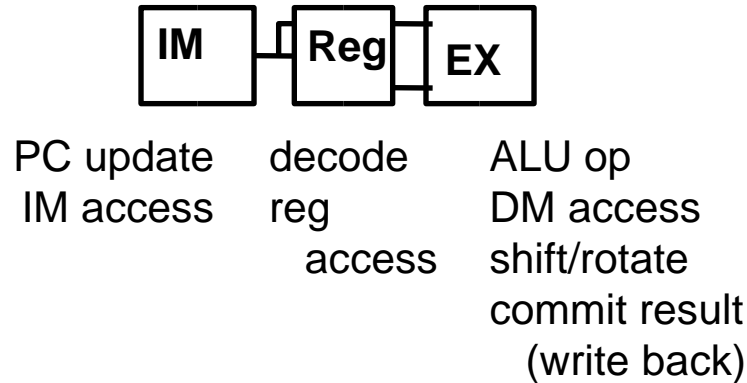
❑ What if the data memory access is twice as slow as the instruction memory?

- make the clock twice as slow or ...
- let data memory access take two cycles (and keep the same clock rate)




Other Sample Pipeline Alternatives

ARM7



XScale



Summary

- ❑ All modern day processors use pipelining
- ❑ Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- ❑ Potential speedup: a CPI of 1 and fast a CC
- ❑ Pipeline rate limited by **slowest** pipeline stage
 - Unbalanced pipe stages makes for inefficiencies
 - The time to “**fill**” pipeline and time to “**drain**” it can impact speedup for deep pipelines and short code runs
- ❑ Must detect and resolve hazards
 - Stalling negatively affects CPI (makes CPI less than the ideal of 1)