


CSL7070: Computer Architecture

Lecture 9, 23rd February 2022

Dip Sankar Banerjee



Indian Institute of Technology, Jodhpur
January-April 2022

CPU Performance Evaluation (CPI)

- tradeoffs
DFS


- Most computers run synchronously utilizing a CPU clock running at a constant clock rate:

Or clock frequency: f

where: Clock rate = $1 / \text{clock cycle}$

$$f = 1 / C$$

(system independent)



- The CPU clock rate depends on the specific CPU organization (design) and hardware implementation technology (VLSI) used.
- A computer machine (ISA) instruction is comprised of a number of elementary or micro operations which vary in number and complexity depending on the the instruction and the exact CPU organization (Design).

- A micro operation is an elementary hardware operation that can be performed during one CPU clock cycle.
- This corresponds to one micro-instruction in microprogrammed CPUs.
- Examples: register operations: shift, load, clear, increment, ALU operations: add, subtract, etc.

CPI → time
IPC → throughput

- Thus: A single machine instruction may take one or more CPU cycles to complete termed as the Cycles Per Instruction (CPI).

// the most imp.

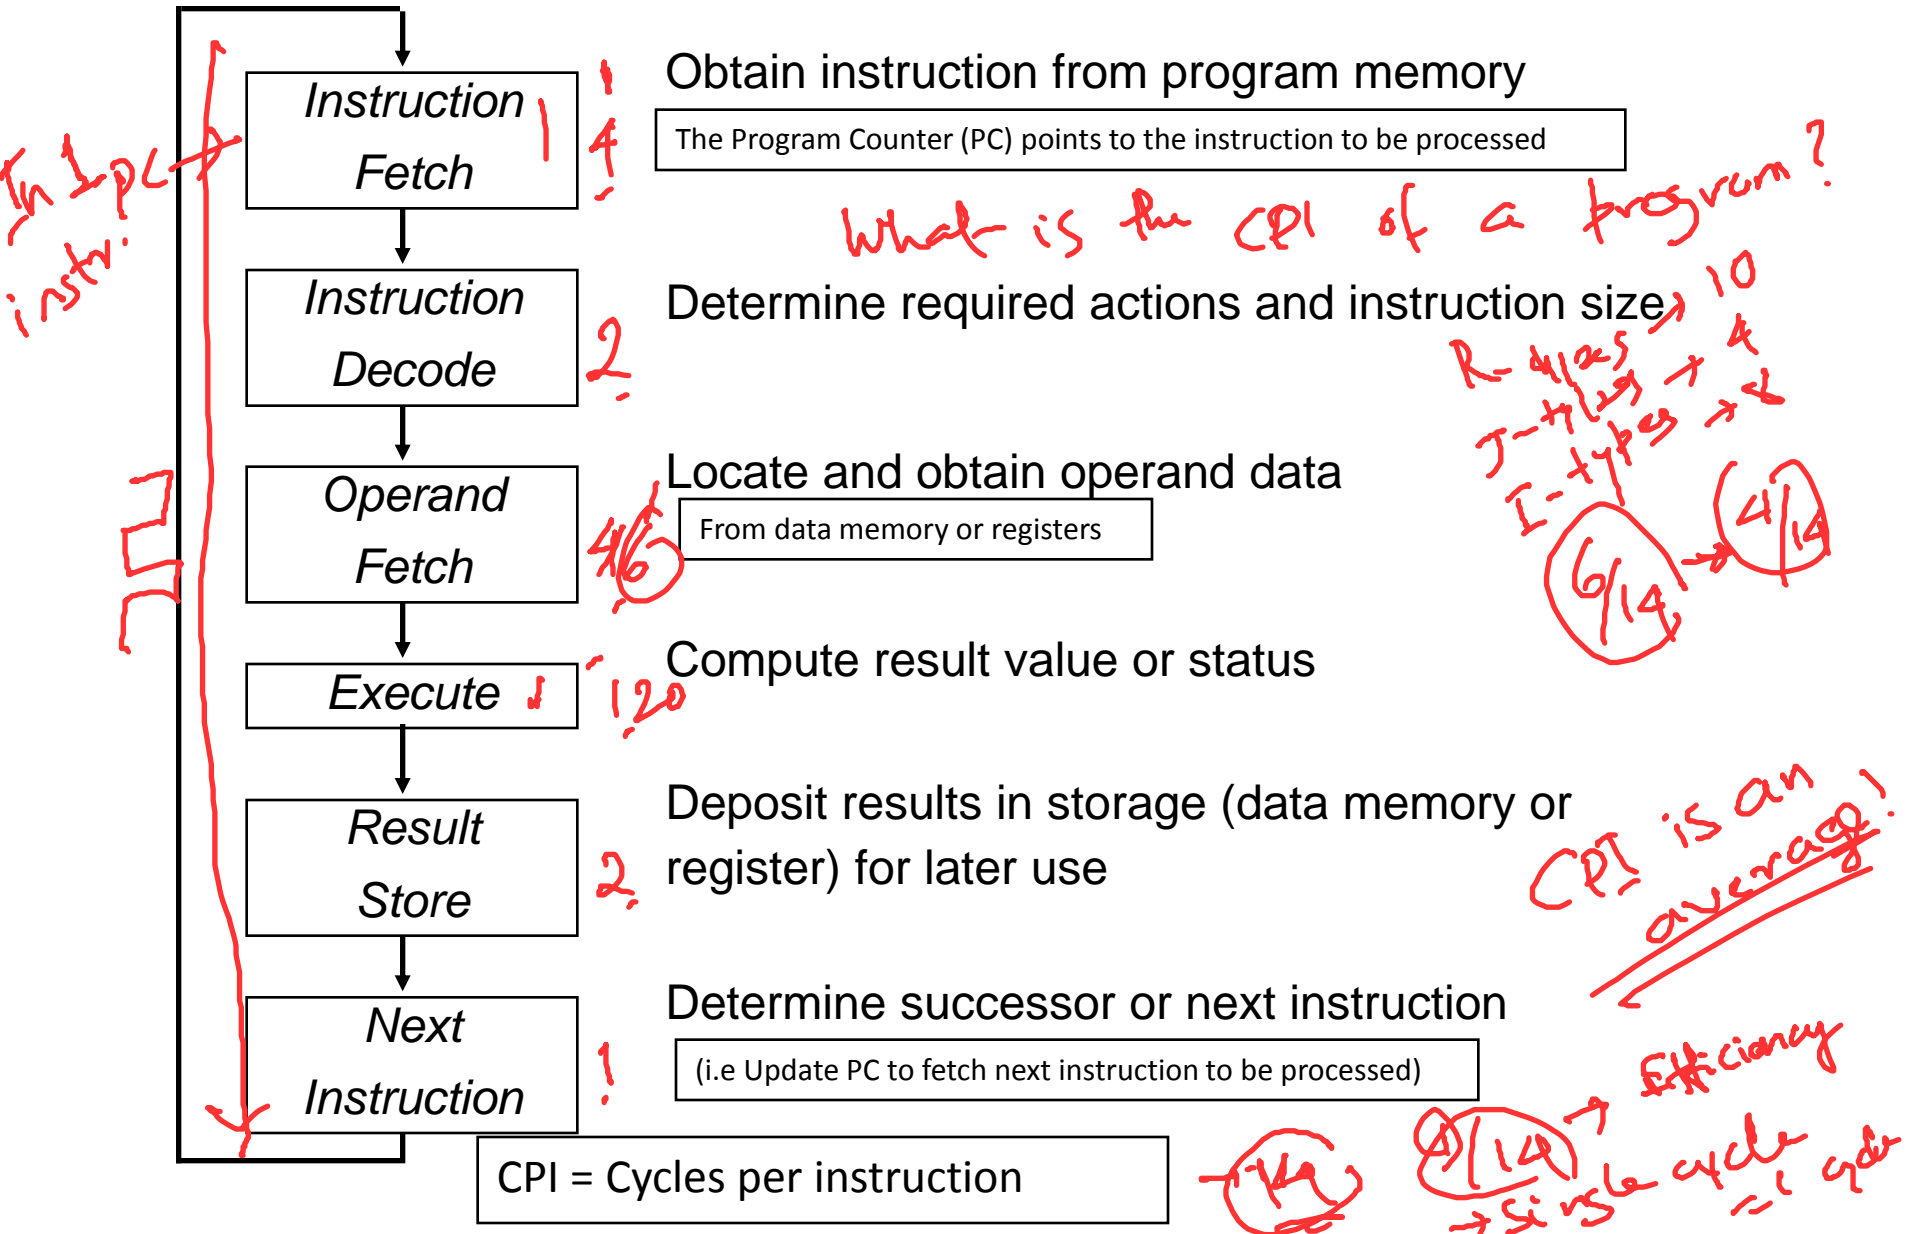
$$\text{Instructions Per Cycle} = \text{IPC} = 1 / \text{CPI}$$

- Average (or effective) CPI of a program: The average CPI of all instructions executed in the program on a given CPU design.

$$\begin{aligned} \text{Cycles/sec} &= \text{Hertz} = \text{Hz} \\ \text{MHz} &= 10^6 \text{ Hz} \quad \text{GHz} = 10^9 \text{ Hz} \end{aligned}$$

CPI ≈ time (synonyms)
6 MHz → $\frac{1}{6 \times 10^6} \times 10^6 \text{ ns} = \frac{1}{6} \text{ ns}$

Generic CPU Machine Instruction Processing Steps



Computer Performance Measures: Program Execution Time

- For a specific program compiled to run on a specific machine (CPU) “A”, has the following parameters:
 - The total executed instruction count of the program. I
 - The average number of cycles per instruction (average CPI). CPI
 - Clock cycle of machine “A” C/CC Or effective CPI
- How can one measure the performance of this machine (CPU) running this program?
 - Intuitively the machine (or CPU) is said to be faster or has better performance running this program if the total execution time is shorter.
 - Thus the inverse of the total measured program execution time is a possible performance measure or metric:

$$\text{Performance}_A = 1 / \text{Execution Time}_A$$

Programs/second Seconds/program

How to compare performance of different machines?

What factors affect performance? How to improve performance?

Comparing Computer Performance Using Execution Time

- To compare the performance of two machines (or CPUs) “A”, “B” running a given specific program:

$$\text{Performance}_A = 1 / \text{Execution Time}_A$$

$$\text{Performance}_B = 1 / \text{Execution Time}_B$$

The two CPUs may target different ISAs provided the program is written in a high level language (HLL)

- Machine A is n times faster than machine B means (or slower? if $n < 1$) :

$$\text{Speedup} = n = \frac{\text{Performance}_A}{\text{Performance}_B} = \frac{\text{Execution Time}_B}{\text{Execution Time}_A}$$

- Example: (i.e Speedup is ratio of performance, no units)

For a given program:

Execution time on machine A: $\text{Execution}_A = 1$ second

Execution time on machine B: $\text{Execution}_B = 10$ seconds

$$\begin{aligned} \text{Speedup} &= \text{Performance}_A / \text{Performance}_B = \text{Execution Time}_B / \text{Execution Time}_A \\ &= 10 / 1 = 10 \end{aligned}$$

The performance of machine A is 10 times the performance of machine B when running this program, or: Machine A is said to be 10 times faster than machine B when running this program.

Proc. A is 10x faster than B
A is 10x faster than B

CPU Execution Time: The CPU Equation

- A program is comprised of a number of instructions executed , I
 - Measured in: instructions/program

AKA Dynamic Executed Instruction Count

- The average instruction executed takes a number of *cycles per instruction (CPI)* to be completed.

- Measured in: cycles/instruction, CPI

Or Instructions Per Cycle (IPC):

$$IPC = 1/CPI$$

- CPU has a fixed clock cycle time C = 1/clock rate

- Measured in: seconds/cycle

$$C = 1/f$$

- CPU execution time is the product of the above three parameters as follows:

$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Executed Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

$$T =$$

execution Time
per program in seconds

$$\cancel{I} \times$$

Number of
instructions executed

$$\text{CPI} \times$$

Average CPI for program

$$C$$

CPU Clock Cycle

Avg.
Total no of cycles
no of instr

(This equation is commonly known as the CPU performance equation)

CPU Average CPI/Execution Time

For a given program executed on a given machine (CPU):

$$\text{CPI} = \frac{\text{Total program execution cycles}}{\text{Instructions count Executed (I)}}$$

(i.e average or effective CPI)

$$\rightarrow \text{CPU clock cycles} = \text{Instruction count (executed, I)} \times \text{CPI}$$

$$\begin{aligned} \text{CPU execution time} &= \\ &= \text{CPU clock cycles} \times \text{Clock cycle} \\ &= \text{Instruction count} \times \text{CPI} \times \text{Clock cycle} \\ T &= I \times \text{CPI} \times C \end{aligned}$$

execution Time
per program in seconds

Number of
instructions executed

Average
or effective
CPI for
program

CPU Clock Cycle

(This equation is commonly known as the CPU performance equation)

CPI = Cycles Per Instruction

CPU Execution Time: Example

- A Program is running on a specific machine (CPU) with the following parameters:

- Total executed instruction count: 10,000,000 instructions
- Average CPI for the program: 2.5 cycles/instruction. i.e 5 nanoseconds
- CPU clock rate: 200 MHz. (clock cycle = $C = 5 \times 10^{-9}$ seconds)

CPU time	=	$\frac{\text{Seconds}}{\text{Program}}$	=	$\frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$
----------	---	---	---	---

- What is the execution time for this program:

$$\begin{aligned}
 \text{CPU time} &= \text{Instruction count} \times \text{CPI} \times \text{Clock cycle} \\
 &= 10,000,000 \times 2.5 \times 1 / \text{clock rate} \\
 &= 10,000,000 \times 2.5 \times 5 \times 10^{-9} \\
 &= 0.125 \text{ seconds}
 \end{aligned}$$

Nanosecond = nsec = ns = 10^{-9} second
MHz = 10^6 Hz

$$T = I \times \text{CPI} \times C$$

✓ Factors Affecting CPU Performance

$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

$$T = I \times \text{Average CPI} \times C$$

	Instruction Count	Cycles per Instruction	Clock Rate (1/C)
<u>Program</u>	✓	✓	✗
<u>Compiler</u>	✓	✗	✗
<u>Instruction Set Architecture (ISA)</u>	✓	✓	✗
<u>Organization (CPU Design)</u>	✗	✓	✗
<u>Technology (VLSI)</u>	✗	✗	✓

$$T = I \times \text{CPI} \times C$$

Aspects of CPU Execution Time

CPU Time = Instruction count executed x CPI x Clock cycle

$$T = I \times \text{CPI} \times C$$

Depends on:
Program Used
Compiler
ISA

Instruction Count I
(executed)

Depends on:
Program Used
Compiler
ISA
CPU Organization

CPI
(Average
CPI)

Clock
Cycle
 C

Depends on:
CPU Organization
Technology (VLSI)

Performance Comparison: Example

- From the previous example: A Program is running on a specific machine (CPU) with the following parameters:

- Total executed instruction count, I: 10,000,000 instructions
- Average CPI for the program: 2.5 cycles/instruction.
- CPU clock rate: 200 MHz.

$$\text{Thus: } C = 1/(200 \times 10^6) = 5 \times 10^{-9} \text{ seconds}$$

- Using the same program with these changes:

- A new compiler used: New executed instruction count, I: 9,500,000

New CPI: 3.0

- Faster CPU implementation: New clock rate = 300 MHz

$$\text{Thus: } C = 1/(300 \times 10^6) = 3.33 \times 10^{-9} \text{ seconds}$$

- What is the speedup with the changes?

$$\text{Speedup} = \frac{\text{Old Execution Time}}{\text{New Execution Time}} = \frac{I_{\text{old}} \times \text{CPI}_{\text{old}}}{I_{\text{new}} \times \text{CPI}_{\text{new}}}$$

$$\begin{aligned} \text{Speedup} &= (10,000,000 \times 2.5 \times 5 \times 10^{-9}) / (9,500,000 \times 3 \times 3.33 \times 10^{-9}) \\ &= .125 / .095 = 1.32 \end{aligned}$$

or 32 % faster after changes.

$$\text{Clock Cycle} = C = 1 / \text{Clock Rate}$$

$$T = I \times \text{CPI} \times C$$

Instruction Types & CPI

- Given a program with n types or classes of instructions executed on a given CPU with the following characteristics:

C_i = Count of instructions of type _{i} executed

CPI_i = Cycles per instruction for type _{i}

e.g ALU, Branch etc.

$i = 1, 2, \dots, n$

Then:

Depends on CPU Design

$$CPI = \text{CPU Clock Cycles} / \text{Instruction Count } I$$

i.e average or effective CPI

Where:

$$CPU \text{ clock cycles} = \sum_{i=1}^n (CPI_i \times C_i)$$

Executed

$$\text{Executed Instruction Count } I = \sum C_i$$

$$T = I \times CPI \times C$$

Instruction Types & CPI: An Example

- An instruction set has three instruction classes:

Instruction class	CPI
A	1
B	2
C	3

e.g ALU, Branch etc. ————

For a specific CPU design

- Two code sequences have the following instruction counts:

Program	Instruction counts for instruction class		
Code Sequence	A	B	C
1	2	1	2
2	4	1	1

- CPU cycles for sequence 1 = $2 \times 1 + 1 \times 2 + 2 \times 3 = 10$ cycles

$$\text{CPI for sequence 1} = \text{clock cycles} / \text{instruction count}$$

$$= 10 / 5 = 2$$

i.e average or effective CPI

- CPU cycles for sequence 2 = $4 \times 1 + 1 \times 2 + 1 \times 3 = 9$ cycles

$$\text{CPI for sequence 2} = 9 / 6 = 1.5$$

$$\text{CPU clock cycles} = \sum_{i=1}^n (CPI_i \times C_i)$$

$$\text{CPI} = \text{CPU Cycles} / I$$

Instruction Frequency & CPI

- Given a program with n types or classes of instructions with the following characteristics:

C_i = Count of instructions of type _{i} executed

$i = 1, 2, \dots, n$

CPI_i = Average cycles per instruction of type _{i}

F_i = Frequency or fraction of instruction type _{i} executed
= $C_i / \text{total executed instruction count} = C_i / I$

Then:

Where: Executed Instruction Count $I = \sum C_i$

$$CPI = \sum_{i=1}^n (CPI_i \times F_i)$$

i.e average or effective CPI

Fraction of total execution time for instructions of type i =

$$\frac{CPI_i \times F_i}{CPI}$$

$$T = I \times CPI \times C$$

Instruction Type Frequency & CPI

Program Profile or Executed Instructions Mix

Base Machine (Reg / Reg)

Depends on CPU Design

$$\frac{CPI_i \times F_i}{CPI}$$

Op	Freq, F_i	CPI_i	$CPI_i \times F_i$	% Time
ALU	50%	1	.5	23% = .5/2.2
Load	20%	5	1.0	45% = 1/2.2
Store	10%	3	.3	14% = .3/2.2
Branch	20%	2	.4	18% = .4/2.2

Given

Typical Mix

Sum = 2.2

$$CPI = \sum_{i=1}^n (CPI_i \times F_i)$$

i.e average or effective CPI

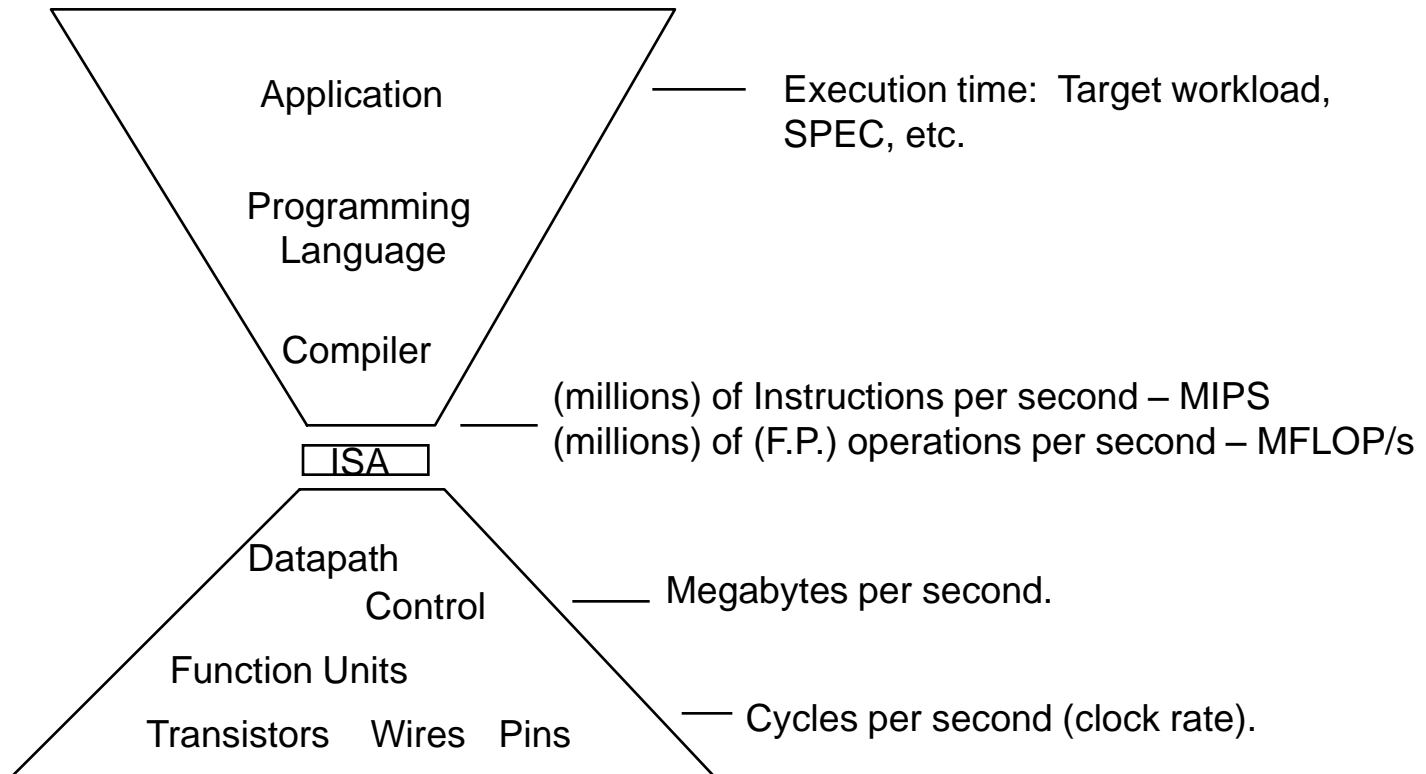
$$CPI = .5 \times 1 + .2 \times 5 + .1 \times 3 + .2 \times 2 = 2.2$$

$$= .5 + 1 + .3 + .4$$

$$T = I \times CPI \times C$$

Metrics of Computer Performance

(Measures)



Each metric has a purpose, and each can be misused.

Choosing Programs To Evaluate Performance

Levels of programs or benchmarks that could be used to evaluate performance:

- Actual Target Workload: Full applications that run on the target machine.
- Real Full Program-based Benchmarks:
 - Select a specific mix or suite of programs that are typical of targeted applications or workload (e.g SPEC95, SPEC CPU2000).
- Small “Kernel” Benchmarks:

Also called synthetic benchmarks

 - Key computationally-intensive pieces extracted from real programs.
 - Examples: Matrix factorization, FFT, tree search, etc.
 - Best used to test specific aspects of the machine.
- Microbenchmarks:
 - Small, specially written programs to isolate a specific aspect of performance characteristics: Processing: integer, floating point, local memory, input/output, etc.

Types of Benchmarks

Pros

- Representative

Actual Target Workload

- Portable.
- Widely used.
- Measurements useful in reality.

Full Application Benchmarks

- Easy to run, early in the design cycle.

Small “Kernel” Benchmarks

- Identify peak performance and potential bottlenecks.

Microbenchmarks

Cons

- Very specific.
- Non-portable.
- Complex: Difficult to run, or measure.

- Less representative than actual workload.

- Easy to “fool” by designing hardware to run them well.

- Peak performance results may be a long way from real application performance

SPEC: System Performance Evaluation Corporation

The most popular and industry-standard set of CPU benchmarks.

Programs application domain: Engineering and scientific computation

- SPECmarks, 1989:
 - 10 programs yielding a single number (“SPECmarks”).
- SPEC92, 1992:
 - SPECint92 (6 integer programs) and SPECfp92 (14 floating point programs).
- SPEC95, 1995:
 - SPECint95 (8 integer programs):
 - go, m88ksim, gcc, compress, li, jpeg, perl, vortex
 - SPECfp95 (10 floating-point intensive programs):
 - tomcatv, swim, su2cor, hydro2d, mgrid, applu, turb3d, apsi, fppp, wave5
 - Performance relative to a Sun SuperSpark I (50 MHz) which is given a score of SPECint95 = SPECfp95 = 1
- SPEC CPU2000, 1999:
 - CINT2000 (11 integer programs). CFP2000 (14 floating-point intensive programs)
 - Performance relative to a Sun Ultra5_10 (300 MHz) which is given a score of SPECint2000 = SPECfp2000 = 100
- SPEC CPU2006, 2006:
 - CINT2006 (12 integer programs). CFP2006 (17 floating-point intensive programs)
 - Performance relative to a Sun Ultra Enterprise 2 workstation with a 296-MHz UltraSPARC II processor which is given a score of SPECint2006 = SPECfp2006 = 1

All based on execution time and give speedup over a reference CPU

Example Problem

Op	Freq	Cycles	CPI(i)	Time
ALU	50%	1	0.5	23%
Load	20%	5	0.5	45%
Store	10%	3	0.3	14%
Branch	20%	2	0.4	18%

Average CPI: 2.2

- How much faster would the machine be if reduce load time to 2 cycle?
- How does it compare with branch prediction to reduce 1 cycle off branch?
- What if 2 ALU instructions can be executed at once?

Example Problem

- Compiler has to decide between two code sequences for a single machine. Based on the hardware implementation there are 3 classes of instructions A, B, C. They require 1, 2 and 3 cycles respectively.
- First code sequence has 5 instr. 2 of A, 1 of B, 2 of C.
- Second has 6 instr. 4 of A, 1 of B, 1 of C
- Which sequence is faster? By how much? What is CPI for each sequence?

Example Problem

- A given application written in Java runs in 15 sec on a machine. A new Java compiler requires only 0.6 as many instr. As old. Unfortunately it raises CPI by 1.1 times.
- How fast can we expect the application to run using this fast compiler?

Example Problem

- Two compilers are being tested for a 4 GHz machine with three different classes of instr. A, B, C requiring 1, 2, 3 cycles respectively. Both compilers are used to produce code for a large software.
- First compiler uses 5 billion Class A, 1 bn class B, and 1 bn Class C.
- Second uses 10 bn A, 1 bn B, 1 bn C.
- Which sequence is faster in mips?
- Which sequence faster according to exec. time?

Example Problem

- A 1 GHz processor takes 100 seconds to execute a program, while consuming 70 W of dynamic power and 30 W of leakage power. Does the program consume less energy in Turbo boost mode when the frequency is increased to 1.2 GHz?

Normal mode energy = $100 \text{ W} \times 100 \text{ s} = 10,000 \text{ J}$

Turbo mode energy = $(70 \times 1.2 + 30) \times 100/1.2 = 9,500 \text{ J}$

Note:

Frequency only impacts dynamic power, not leakage power. We assume that the program's CPI is unchanged when frequency is changed, i.e., exec time varies linearly with cycle time.